

A NEW RENDERING AND PLUME SIMULATION FRAMEWORK FOR VISUALIZING TAAL VOLCANO ERUPTION PLUMES

A Thesis Proposal

presented to

the Department of Software Technology
College of Computer Studies
De La Salle University

In partial fulfillment
of the requirements for the degree of

Bachelor of Science in Computer Science
Major in Software Technology

by

BERNANDINO, Adrian Rafael
BORRAMEO, Anton Miguel
GALAN, Russell Emmanuel
GANAYO, Reinman Geoffrey

Neil Patrick Del Gallego, Ph.D.
Adviser

August 1, 2026

Abstract

The Taal Volcano in the Philippines remains one of the country's most active volcanoes due to its recurring eruptive activity, highlighting the continued need for effective tools that support hazard communication, disaster preparedness, and scientific visualization. Interactive three-dimensional computer graphics simulations have demonstrated their potential to provide intuitive visual representations of volcanic eruption plumes while allowing users to explore eruption scenarios in an immersive virtual environment. AnitoPlume, an existing graphics-based Taal Volcano plume simulator, has successfully demonstrated real-time visualization capabilities using a particle-based simulation model; however, it relies on legacy graphics rendering technologies and software architectures that limit extensibility, maintainability, and the efficient utilization of modern graphics hardware. This study proposes a modern 3D graphics rendering framework for interactive Taal Volcano plume simulation by porting and extending AnitoPlume, redesigning its rendering architecture while preserving its established particle-based simulation model and core visualization techniques. The proposed framework aims to modernize the rendering pipeline, graphics resource management, shader implementation, and software architecture to provide a more scalable, maintainable, and extensible platform for future graphics-based scientific visualization applications. A design and development research methodology will be employed to reimplement AnitoPlume using a contemporary graphics rendering framework and evaluate its rendering performance, graphics resource utilization, scalability, and visual fidelity through quantitative performance metrics and qualitative visual assessment.

Keywords: Computer graphics, 3D graphics, Graphics rendering, Volcanic plume simulation, Taal Volcano, Particle systems, AnitoPlume

Chapter 1

Introduction

1.1 Background of the Study

Volcanic eruptions are among the most destructive natural hazards, producing ash plumes, pyroclastic flows, and volcanic gases that threaten nearby communities and critical infrastructure. Effective visualization of these hazards plays an important role in disaster preparedness, scientific communication, and public education. Traditionally, volcanic simulations have relied on numerical models that prioritize scientific accuracy for hazard prediction and analysis. While these models provide valuable quantitative information, they are often computationally intensive and lack the interactive capabilities necessary for real-time visualization (Rossant & Rougier, 2021).

The original AnitoPlume was developed as a real-time 3D simulator for Taal Volcano. Replicating the 2020 eruption and providing an interactive, and immersive, application for use in disaster communication and awareness (Del Gallego et al., 2026). The application can be broken down into its three major components: a terrain model, the plume simulations, and the UI. The terrain component renders a detailed elevation model of Taal Volcano, with collisions and maps, serving as the basis for plume behavior regarding terrain. The plumes use a Lagrangian, particle-based approach to represent eruption dynamics through smoke layers, free spheres, falling spheres, and stagnation states. These allow the plume to rise, spread, collapse, and interact with wind and terrain. The user interface provides control over eruption parameters, wind settings, camera and display modes, etc.




Figure 1: Terrain models of Taal Volcano inside the AnitoPlume and different plume simulations of AnitoPlume. (Del Gallego et al., 2026)




Figure 2: User interface of AnitoPlume. The Taal Volcano model is rendered in the center of the simulator. The photo reference is on the left. (Del Gallego et al., 2026)

Graphics Application Programming Interfaces (APIs) define how applications communicate rendering operations to the graphics processing unit (GPU). Unlike traditional implicit graphics APIs, modern explicit graphics APIs require applications to explicitly manage rendering state, resource allocation, and command submission, thereby reducing driver overhead and improving utilization of modern multicore processors and graphics hardware (Ping et al., 2026).

Despite the theoretical performance advantages of modern rendering pipelines, a significant challenge remains: the architectural enhancement of established scientific tools. Therefore, investigating how GPU control impacts rendering throughput, memory bandwidth utilization, and overall visual fidelity provides essential insights for the future development of highly scalable scientific visualization systems (Unterguggenberger, Kerbl, & Wimmer, 2022). While AnitoPlume demonstrates that a volcanic plume can be visualized interactively at real-time framerates, its existing graphics architecture limits the possible terrain and plume optimizations.



Figure 3: Overview of older and modern graphics APIs. From top to bottom, left to right: OpenGL ES, OpenGL, DirectX 11, MetalFX, Vulkan, DirectX 12. (TechArx Gaming Staff, 2016)

The proposed AnitoPlume V2 therefore aims to port the visual system to a new rendering framework, and revisit the existing plume simulation logic. Improving on, but not limited to: GPU control and utilization, shading flexibility, rendering efficiency, and simulation logic improvements. In doing so, the project will improve how the terrain and plumes are rendered, with emphasis on higher fidelity terrain models and plume simulations at least as detailed as the current representation.

1.2 Research Objectives

This study aims to port and extend AnitoPlume to a modern graphics framework, preserving its established particle-based simulation model, evaluate the resulting

improvements, and assess whether the system maintains an interactive and usable experience.

The specific objectives of this research are as follows:

1. To benchmark GPU resource synchronization and cross-API shader translation for AnitoPlume’s volumetric plume data, establishing a baseline for porting the simulator to the new rendering framework;
2. To redesign AnitoPlume’s software architecture by migrating to alternative rendering frameworks that support explicit-based architectures, physically-based lighting (PBL), such as but not limited to Vulkan and DirectX 12;
3. To analyze the rendering efficiency of the modernized pipeline against the previous implementation using GPU profiling metrics, such as, but not limited to, frame pacing, shader execution time, memory bandwidth utilization, and overall render throughput, under identical hardware constraints; and
4. To validate the overall user experience of the modernized simulator by conducting usability testing with target beneficiaries.

1.3 Scope and Limitations of the Research

This study is limited to porting and extending AnitoPlume, an existing Taal Volcano plume visualization system, from its current graphics API (OpenGL) to a modern graphics API, such as but not limited to Vulkan or DirectX 12. In contrast to legacy APIs, where the graphics driver automatically manages GPU state, memory allocation, and synchronization on the application’s behalf, APIs require the application to directly define and control these aspects, including Pipeline State Objects (PSOs), command lists, and GPU resource synchronization. This shift is the basis for the “explicit” rendering architecture investigated in this research. The study specifically examines pipeline state management, command recording, shader portability across graphics APIs, and GPU resource synchronization, while strictly preserving AnitoPlume’s existing particle-based plume simulation model. Modifications to the underlying scientific simulation algorithms, volcanic plume physics, or eruption modeling parameters are outside the scope of this work. Furthermore, the study is scoped to graphics API-level constructs rather than the abstractions of any single third-party rendering framework or engine, ensuring that the resulting architecture and findings remain applicable across explicit-API implementations rather than being dependent on one specific framework.

The usability testing phase will involve 10 to 20 voluntary participants, recruited according to a tiered priority order: PHIVOLCS personnel will be prioritized as the primary participant group given their domain expertise in volcanology; should this pool be insufficient, recruitment will expand, in order, to individuals who witnessed a Taal Volcano eruption firsthand, then to volcano/geohazard simulation enthusiasts, and lastly to computer science students or developers with a relevant background in 3D graphics or software development. For PHIVOLCS personnel, recruitment will be initiated via email correspondence with the agency; if permission is granted, in-person coordination will also be conducted. For the remaining participant categories, recruitment will be conducted through a general online survey containing screening questions to confirm which category, if any, each respondent qualifies under. Data collection will take place at the DLSU GameLab, Gokongwei Building, Room 411 (G411). After completing a standardized set of interaction tasks with the ported simulator, participants will complete an anonymous, self-administered questionnaire using the System Usability Scale (SUS) to evaluate application responsiveness, stability, and visual parity relative to the original AnitoPlume system. This phase is expected to take 2 to 4 weeks within the systems evaluation timeframe of the study.

This research is limited to a single case study, the AnitoPlume Taal Volcano plume simulator, and its findings on rendering performance and usability are not guaranteed to generalize to other volcanic visualization systems or simulation platforms. Benchmarking results are further constrained to the specific hardware configurations used during testing and may not be representative of performance across all GPU vendors or consumer hardware. Finally, given the sample size of 10 to 20 participants, whose composition may range from an entirely PHIVOLCS-based sample to a mix across the tiered participant categories depending on PHIVOLCS availability, usability findings are intended to provide indicative, rather than statistically generalizable, insight into user experience.

1.4 Significance of the Research

This study contributes to both computer graphics research and scientific visualization by providing empirical evidence on the application of modern graphics architectures to real-time volumetric scientific visualization, while also offering immediate, practical value to Philippine disaster risk reduction efforts.

For Philippine communities and disaster management stakeholders, this study offers direct, immediate benefits. Taal Volcano remains under close monitoring due to its recurring eruptive activity, yet accessible, interactive 3D tools for com-

municating volcanic hazards to the public and to local decision-makers remain limited. By modernizing and extending AnitoPlume into a more performant, maintainable, and scalable system, this study lays the groundwork for AnitoPlume V2 to potentially serve as one of the first accessible 3D visualization tools for volcanic hazard communication in the Philippines. Such a tool could support disaster preparedness education, hazard awareness campaigns, and public information efforts led by local government units (LGUs) and agencies such as PHIVOLCS, by allowing communities and responders to visualize eruption scenarios in a more intuitive and immersive manner than static hazard maps alone. While this study does not directly evaluate deployment with disaster response agencies or affected communities, it establishes the technical foundation necessary to pursue such applications in future work.

For researchers in computer graphics, the study offers a detailed case study demonstrating the architectural modifications required when transitioning between two fundamentally different graphics APIs.

For developers of scientific visualization software, the findings may serve as a reference for future projects that require designing high-performance rendering systems for similarly data-intensive, real-time visualization applications.

For future researchers, this work establishes a scalable rendering foundation that supports integrating advanced graphics techniques, including hardware-accelerated ray tracing, compute shaders, GPU-driven rendering, and other contemporary visualization methods. Consequently, the study extends the long-term applicability of the original AnitoPlume project beyond its initial legacy implementation, while positioning it as a potential stepping stone toward practical, real-world disaster communication tools for volcanic hazards in the Philippines.

Chapter 2

Related Works

2.1 Taal Volcano Primer




Figure 4: From left to right, top to bottom: Taal Volcano during 2015, 2019, 2023, and 2026. (TechArx Gaming Staff, 2016; Flores, 2019; Daguno-Bersamina, 2023; Mallari Jr., 2026)

Taal Volcano is a complex volcano located in southwestern Luzon, Philippines, situated in the middle of Taal Lake (Smithsonian Institution, 2026). Figure 4 illustrates the volcano across different time periods, including its primary eruption

phases. As one of the most active volcanoes in the country, the Philippine Institute of Volcanology and Seismology (PHIVOLCS) has recorded 36 historical eruptions at Taal (PHIVOLCS, 2024). Its frequent activity poses a major threat to local communities and the densely populated Metro Manila area. Although official hazard maps designate only the volcano island as a Permanent Danger Zone (PDZ) and outline a high-risk perimeter of 15 km, Taal’s actual danger extends much further due to the regular release of toxic gases. Consequently, Metro Manila remains highly vulnerable to these emissions despite being located approximately 81 km away from the volcanic center.

This vulnerability was clearly demonstrated during the January 12, 2020 eruption, when a phreatomagmatic event occurred at the main crater, producing a massive column of ash and steam. The severity of the eruption led PHIVOLCS to raise the alert status to Level 4, indicating that a hazardous eruption was imminent. The event was classified as a Volcanic Explosivity Index (VEI) 4 eruption, depositing an estimated tephra fallout volume of 0.097 km^3 (Balangue-Tarriela et al., 2022). Due to the immediate risks, a mass evacuation of roughly 376,000 people was carried out in the surrounding areas, including the provinces of Batangas and Cavite (Renzo, 2020; Del Castillo, Paraiso, Vicente, Jamero, & Narisma, 2020).

The 2020 event is the only major eruption of Taal to be thoroughly documented, thanks to modern communication and technologies like accessible digital cameras, crowdsourcing, and high-resolution seismic sensors. This availability of real eruption footage and geographical data provides a strong foundation for developing visual simulations of the volcano.

2.2 Geoscientific Modeling and Volcanic Crisis Communication

2.2.1 Limitations of Current Numerical Models

In volcanology, the standard approach for evaluating risk and hazard zones relies on numerical simulations that model the physics of ash transport and tephra fallout. Volcanologists widely use Eulerian models such as FALL3D (Folch, Costa, & Macedonio, 2009), Tephra2 (Bonadonna, 2006; Connor & Connor, 2006), and TephraProb (Biass, Bonadonna, Connor, & Connor, 2016), alongside Lagrangian models like HYSPLIT (Stein et al., 2015), for scientific analysis and hazard mapping. These models are highly valued for their quantitative accuracy and their

capacity to incorporate atmospheric variables.

However, these simulations are built for expert analysis rather than public situational awareness, requiring high-performance computing and substantial processing time. While these methods are essential for long-term hazard forecasting, they lack the interactive flexibility and 3D immersion needed for real-time visualization applications. To address this, our proposed software, AnitoPlume, serves as a computer graphics application that complements these numerical models by offering an interactive testbed for 3D volcanic eruption visualizations of Taal Volcano.

2.2.2 Advancements in 3D Hazard Visualization

Traditionally, volcanic hazards and crisis information are communicated using 2D maps that display contour lines and heat maps for ash distribution, pyroclastic flow boundaries, and tephra thickness (Nave, Isaia, Vilardo, & Barclay, 2010; Thompson, Lindsay, & Gaillard, 2015; Fearnley & Beaven, 2018; Esposti Ongaro, Cerminara, Charbonnier, Lube, & Valentine, 2020). Although 2D representations remain the standard for scientific experts, research indicates that integrating "spatial familiarity" is critical for prompting effective public responses (Preppernau & Jenny, 2015). Comparative studies further demonstrate that general users often prefer 3D perspectives over 2D maps, as 3D views make it easier to understand complex terrain features and identify viable evacuation routes (Thompson, Lindsay, & Leonard, 2017; Driedger et al., 2024).

To improve public communication, several research initiatives have explored immersive and interactive applications using computer graphics (Zhang, Kong, Li, Wang, & Qin, 2017; Lastic et al., 2022; Pretorius et al., 2024), virtual reality (VR) (Tibaldi et al., 2020; Asgary, Bonadonna, & Frischknecht, 2019; Li, 2023), or gamification (Kerlow, Pedreros, & Albert, 2020). Tools such as VolcanoVR (Hansen et al., 2025), VRVOLC (Delage et al., 2021), and HoloVulcano (Asgary et al., 2019) utilize commercial game engines to let stakeholders view complex volcanic phenomena in safe, immersive environments.

Distinct from these commercial engine approaches, the original AnitoPlume framework was developed using a custom, open-source stack (C++, OpenGL, Eigen, dearIMGUI) to bypass general-purpose engine abstractions and optimize real-time rendering performance. Furthermore, while generic eruption simulations exist (Lastic et al., 2022; Pretorius et al., 2024), AnitoPlume specifically targeted the high-risk Taal "Decade Volcano" by integrating LiDAR and OpenTopography terrain datasets with real-time plume dynamics. Building directly upon this

specific, site-tailored approach, our work modernizes this legacy architecture by replacing its original rendering backend with a next-generation graphics API, significantly enhancing visual fidelity and pipeline efficiency.

2.3 Geosciences and computer graphics

Volcanic modeling paradigms in geosciences and computer graphics differ significantly in both methodology and primary intent. Within geoscientific and volcanological research, the priority is predictive accuracy and scientific forecasting. In this context, numerical solvers are essential for quantifying physical parameters such as mass eruption rates and tephra dispersal (Kavanagh, Engwell, & Martin, 2018; Poland & Anderson, 2020; Esposti Ongaro et al., 2020).

Conversely, in the field of computer graphics, volcanic modeling is treated as a subfield of real-time rendering and procedural simulation, focusing primarily on particle systems and fluid dynamics (Stora, Agliati, Cani, Neyret, & Gascuel, 1999; Roth & Guritz, 1995; Lastic et al., 2022; Li, 2023; Pretorius et al., 2024). Research in this domain emphasizes real-time user interactivity and immersion, meaning that visual effects are often approximations and physical models are simplified. Consequently, these techniques prioritize visual realism and low-latency performance over the complex computations found in traditional numerical solvers.

2.4 Particle systems

2.4.1 Real-time particle systems

Within computer graphics applications, volcanic plumes can be simulated utilizing particle systems, an approach similar to the representation of other atmospheric phenomena such as smoke, fog, and clouds (Rautenhaus et al., 2017; Engel et al., 2004; Akenine-Moller et al., 2019). Real-time (RT) particle systems model these phenomena as collections of "fuzzy objects" (Reeves, 1998) that lack well-defined surface, distinguishing them from traditional 3D models composed of explicit polygonal surfaces. Every constituent particle is dynamic, continuously altering its morphology and position over a discrete lifespan that spans from its birth to its eventual dissipation.

Modern frameworks replicate the stochastic properties of individual particles

by exposing various parameters that users can adjust to mimic specific visual effects like fire, smoke, or fog. This methodology is formally classified as a Lagrangian perspective. Particles originate from an emitter, which serves as the spawning source within a 3D grid. The geometric configuration of the emitter plays a critical role in shaping the final visual output; for instance, a conical emitter is highly suitable for volcanic plume simulations because it mirrors the upward-expanding, conical trajectory characteristic of eruption columns.

Volumetric density is typically regulated via an "emission rate" parameter, while lifetime behaviors are dictated by variables governing velocity, color, and transparency transitions, or by enabling particles to dissolve upon intersecting solid geometry. Rendering is generally executed using simple primitives, such as points or lines, or via screen-aligned 2D textures known as billboards. In our implementation, we developed AnitoPlume to sample dynamically from a curated five-set texture library of smoke assets engineered to closely replicate the visual characteristics of an authentic volcanic plume.

Although RT particle systems are computationally efficient and highly capable of approximating plume behavior, they demand extensive manual parameter tuning (e.g., managing dimensions, emission frequencies, lifespans, velocities, and color gradients). Furthermore, they struggle to natively model key physical behaviors, such as the internal density distribution across the particle volume. Consequently, incorporating physical fluid dynamics becomes necessary to capture authentic plume mechanics (Section 2.5.2).

2.4.2 Particle systems and fluid dynamics

While real-time particle frameworks offer efficient approximations and excel at high-frame-rate visualization, they fundamentally lack the mechanism to resolve complex physical fluid interactions (Nishita & Dobashi, 2001). A volcanic plume represents a complex, multi-phase mixture of gases and ash; standard real-time particle systems cannot autonomously compute environmental interactions involving atmospheric clouds, wind fields, air entrainment, and gravitational forces (Lastic et al., 2022; Pretorius et al., 2024). To resolve this limitation, several methodologies integrate fluid solvers to govern underlying particle motion, which is subsequently visualized via a real-time engine or exported to offline production suites like Houdini (Fedkiw, 2003; Pfaff, Thurey, & Gross, 2012; Lastic et al., 2022; Pretorius et al., 2024).

Purely Lagrangian frameworks utilized in RT graphics incur steep computational overhead when managing massive particle counts ($> 500K$ on mid-range

graphics hardware), rendering them impractical for highly interactive, real-time applications. To address this, researchers leverage Eulerian architectures, where grid-based solvers prove more efficient for representing continuous media like fluids and capturing high-fidelity scientific details (Jänicke & Scheuermann, 2010; Pfaff et al., 2012; Saidi, Rismanian, Monjezi, Zendehbad, & Fatehiboroujeni, 2014). Hybrid Eulerian–Lagrangian systems have also been introduced to exploit the dual advantages of both paradigms, with recent implementations deployed in cloud and cyclone modeling (Vimont et al., 2020; Stam, 2023; Pretorius et al., 2024; Amador Herrera et al., 2024). To contextualize this, Figure 5 outlines different kinds of particle-based modeling:




Figure 5: Overview of different particle systems and fluid simulations. From top to bottom, left to right: graph neural network (GNN) for physics simulations, procedural particle system for rendering fuzzy physical phenomena, high-resolution time evolution of a smoke explosion, compiled different particle simulations. (Sanchez-Gonzalez et al., 2020; Reeves, 1998; Selle et al., 2005; Nowrouzezahrai, 2006)

The core simulation pipeline of AnitoPlume closely mirrors the Lagrangian volcanic plume model developed by Lastic et al.. Within this framework, column dynamics are resolved via a top-down architectural approach: an ascending circular cross-section simulates the rising eruption column emerging from the volcanic crater, which subsequently spawns rotating spheres for each discrete layer. This rotational behavior explicitly accounts for the localized vorticity and plume rotation documented in physical volcanic events (Nabizadeh, Roy-Chowdhury, Yin, Ramamoorthi, & Chern, 2024; Braun, Bender, & Thuerey, 2025). In modern-

izing the application’s graphics API, these rotating spheres are translated into instanced, screen-aligned billboard textures, which are dynamically sampled from the system’s five-set smoke asset library based on the particle’s instantaneous altitude.

2.5 Terrain rendering

Virtual terrains are typically represented through digital elevation models (DEM), such as textured meshes or heightmaps. Approaches to constructing these terrains can generally be categorized by how the underlying elevation data is obtained. *Acquisition-based techniques* derive terrain data from real-world sources, such as satellite imagery (e.g., Google Maps/Earth) or LiDAR scanning, allowing the reconstruction of terrain that reflects actual geographical landscapes. In contrast, *procedural techniques* generate terrain algorithmically, typically through noise-based functions, without reliance on real-world source data. Regardless of acquisition method, elevation data is commonly represented as a *heightmap*, a widely adopted format consisting of a 2D raster grid of elevation values, which is utilized to extrude a flat plane into a three-dimensional surface.

When modeling terrains inspired by actual geographical landscapes, data must be sourced from geographic information systems, including satellite imagery or LiDAR maps. Aside from ground elevation, raw LiDAR scans also include reflections from vegetation, structures, and other above-ground features that must be filtered out before a usable terrain surface can be extracted. Figure 6 illustrates this filtering process, showing the progression from raw above-ground height measurements to an isolated Bare Earth DEM suitable for terrain extrusion.




Figure 6: Example of processed LiDAR output layers. From (U.S. Geological Survey, 2018).

While converting satellite and LiDAR data into a heightmap to generate a 3D terrain within visualization software is a direct methodology, it remains susceptible to data gaps and structural inaccuracies, particularly when the underlying DEMs are captured by low-resolution sensors.

Prior works focus on the volcanic eruption behaviors, such as lava propagation or plume dynamics, while rarely accounting for volcanic topography during the simulation process. Because AnitoPlume was designed as a holistic 3D simulator of the Taal Volcano, terrain rendering constitutes a core component of its architectural framework. The baseline 3D model of the Taal Volcano Island was constructed using an acquisition-based approach: LiDAR data was extracted from opentopography.org (openTopography, 2025) and processed within QGIS (Moyroud & Portet, 2018), a geographic information system utilized to generate the requisite heightmap. This heightmap was leveraged to produce an initial 3D mesh of Taal, which was subsequently refined using Blender (Blender Foundation, 2025). Any incomplete segments within the original LiDAR data were manually resolved using digital inpainting, sculpting, and smoothing workflows. Figure 7 summarizes this end-to-end pipeline, from LiDAR-derived heightmap conversion

in QGIS through mesh refinement in Blender.




Figure 7: AnitoPlume’s terrain modeling workflow of Taal Volcano. From (Del Gallego et al., 2026)

2.6 Volcanic simulations

Some computer graphics researchers focus on the generation of visually convincing lava flow animations. For instance, Zhang et al. created a hybrid framework engineered to simulate lava viscosity while accounting for thermal transfer and surface friction mechanics. To capture granular-fluid flows—which are highly applicable to mass-wasting events such as landslides, avalanches, and lava progression—researchers frequently adopt debris flow models because they synthesize foundational concepts from fluid, solid, and soil mechanics (Iverson & George, 2014; George & Iverson, 2014; Trujillo-Vela, Ramos-Cañon, Escobar-Vargas, & Galindo-Torres, 2022).

As for the simulation of volcanic plumes, it typically employs the use of atmospheric transport or scattering models, ensuring that the simulated plume interacts realistically with the ambient air or wind (Lastic et al., 2022; Pretorius et al., 2024). New novel fluid such as Coadjoint Orbit FLIP (CoFLIP) (Nabizadeh et al., 2024) and Adaptive-FLIP (Braun et al., 2025), both of which operate as modern extensions of fluid implicit particle (FLIP) formulation (Brackbill & Ruppel, 1986). Among these developments, the framework proposed by Lastic et al. represents the first work observed so far that can perform real-time plume simu-

lation, due to utilizing a streamlined Lagrangian simulation. Building upon this, Pretorius et al. proposed a coupling model of explosive eruptive behaviors directly with atmospheric simulations to render volcanic skiescapes.

Figure 8 illustrates representative outputs from real-time plume simulation frameworks, contrasting the Lagrangian-based approach of Lastic et al. and the coupled skyscape model of Pretorius et al. with AnitoPlume’s own real-time volcanic plume simulation of Taal Volcano, which serves as the baseline system extended in this work.




Figure 8: Volcanic simulations from (Del Gallego et al., 2026) (top), (Lastic et al., 2022) (bottom-left), and (Pretorius et al., 2024) (bottom-right).

Alternative domains of research prioritize the analysis and interpretation of eruption data (Pardini et al., 2024). Recently, several works have expanded to include deep learning and semantic segmentation techniques designed to automatically detect and isolate volcanic plumes within optical camera imagery (Wilkes, Pering, & McGonigle, 2022; Guerrero Tello, Coltelli, Marsella, Celauro, & Palenzuela Baena, 2022; Torrisi, Corradino, Cariello, & Del Negro, 2024; Swetha, Datla, & Vishnu, 2023).

2.7 AnitoPlume

AnitoPlume is the first interactive, real-time 3D graphics-based simulation application explicitly tailored for visualizing the changing landscape of the Taal Volcano island in the Philippines and its corresponding eruption plumes. Developed primarily to address the challenges of volcanic hazard and crisis communication, the simulator serves as an immersive, lightweight alternative to traditional 2D heatmaps and highly complex, non-interactive geoscientific numerical solvers (such as FALL3D and Tephra2). While geoscientific solvers excel at quantitative predictive accuracy, they require heavy high-performance computing resources and significant processing time, making them unviable for applications requiring immediate visual feedback, interactive flexibility, and low-latency user immersion. AnitoPlume fills this gap by merging computer graphics techniques with site-specific real-world terrain details to facilitate risk awareness and “what-if” situational modeling (Del Gallego et al., 2026). To visually demonstrate these differences in rendering approach and graphical fidelity, Table 1 compares the output of previous plume simulations with the real-time terrain integration of our proposed simulator.

	Geoscientific Models (e.g., FALL3D, Tephra2)	Computer Graphics Related Work (e.g., Lastic et al. (2022))	AnitoPlume (Proposed Simulator)
Method	Numerical solvers, Eulerian atmospheric dispersion models, Tephra accumulation and dispersion.	Eulerian, Lagrangian, or hybrid particle systems approximation.	Lagrangian system with synchronized multi-vent plume logic.
Visualization	Static 2D/3D low-fidelity contours, graph-based and/or scientific visualization, heatmaps.	3D scene using simple spheres, billboard particles, static planes, or basic landscapes.	3D scene using LiDAR and OpenTopography terrain data, dynamic textures, textured particles, interactive UI with plume trackers.
Performance	High accuracy at the cost of high computation time.	Real-time performance prioritized for visual realism. Approximation of physical eruption behavior.	Real-time performance with improved visual realism. Optimized for 30 FPS. E.g., Terrain-specific model through multi-source 3D modeling workflow.
Temporal Representation	Focuses on snapshots of single eruptive events.	Continuous loop simulations with emphasis on plume structure alone.	Dynamic textures: Multi-year landscape visualization (2015–2023), different plume particles based on duration.
User Interface (UI)	Displays comprehensive outputs and calculations. E.g., dispersal, mass accumulation, probability maps.	Parameter tuning for visual and physics adjustments, either via the UI or via text/command-line input.	A user application requiring no text prompt. Revised layout for parameter tuning. UI is suited for “what-if” scenarios because of integrated hazard trackers.
Validation Strategy	Field-constrained analysis of tephra fall deposits.	Qualitative visual comparisons	Quantitative similarity analysis — structural image similarity (SSIM) and root mean squared error (RMSE). Usability testing conducted.

Table 1: Comparison of previous plume simulation (left) implemented by Lastic et al. (2022) and geoscientific numerical solvers (middle) with the AnitoPlume simulator (right) by Del Gallego et al. (2026).

2.7.1 Graphics Architecture and System Performance

To bypass the overhead, heavy abstractions, and unneeded general-purpose features of commercial game engines, AnitoPlume was custom-built using a lean, native C++ open-source API stack consisting of OpenGL for graphics rendering, Eigen for numerical linear algebra, and dearIMGUI for the user interface layout. This custom application architecture optimizes real-time performance on consumer-grade hardware. By keeping the system dependencies minimal, the simulator successfully achieves real-time interactivity, running at an average rendering rate of 30 frames per second (FPS) on a system with 6 GB of video RAM (VRAM) capped at 2080 x 2080 resolution (Del Gallego et al., 2026).

2.7.2 Usability Design and Interactive Features

The simulator also introduced several usability improvements over previous volcanic visualization systems, including simplified camera controls, interactive parameter editing, real-time plume direction tracking, and contextual information about nearby communities affected by volcanic activity. These additions improved both accessibility and user experience while preserving interactive performance (Del Gallego et al., 2026).

2.7.3 System Validation and Evaluation Metrics

To evaluate the effectiveness of the simulator, Del Gallego et al. conducted quantitative and qualitative assessments. Visual fidelity was measured using the Structural Similarity Index (SSIM) and Root Mean Squared Error (RMSE) by comparing simulated eruption images with real-world photographs and surveillance footage of Taal Volcano. The study reported high visual similarity, indicating that the synthesized plume closely resembled actual volcanic eruptions. In addition, usability testing using the System Usability Scale (SUS) demonstrated that participants generally found the simulator easy to use and suitable for interactive exploration.

2.7.4 Architectural Limitations and Future Trajectories

Despite these accomplishments, the primary contribution of AnitoPlume lies in the design of its visualization framework rather than the underlying graphics API.

The study focused on plume simulation, terrain modeling, visual fidelity, and usability evaluation, while the OpenGL implementation served primarily as the rendering platform supporting these objectives. Consequently, the original work did not investigate how the rendering architecture itself might influence performance, resource utilization, maintainability, or future extensibility.

The authors also acknowledged several opportunities for future development, which include:

1. Photorealistic Rendering
2. Expanded Physics-Based Mechanics
3. Image-Based Parameter Estimation
4. Geoscience Data Integration
5. Automated Terrain Synthesis

These recommendations highlight the critical need for a modern rendering architecture and a low-level API like DirectX 12, which are capable of supporting such intensive graphics and parallel computing workloads while maintaining real-time interactive performance.

2.8 API Migration in Rendering Systems

The migration of rendering software between graphics APIs is a well-known challenge in computer graphics. It often involves more than just replacing function calls. A successful port must account for differences in shader languages, resource-binding models, render-target handling, memory synchronization, and execution semantics. In practice, this means that a program cannot simply be translated line by line; instead, the rendering architecture must be reinterpreted in the context of the target API.

Previous work in graphics engineering has shown that API migration is most successful when the original system is structured in a modular way. Systems that separate rendering logic, scene management, and platform-specific code are often easier to port than tightly coupled implementations. This is particularly relevant for research software, where maintainability and clarity are essential. The process of porting from OpenGL to DirectX, therefore, becomes not only a technical exercise, but also an opportunity to improve the software architecture.

2.8.1 Cross-API Shader Translation

A major issue in porting visual effects from OpenGL to DirectX 12 is shader translation. OpenGL workflows usually rely on the OpenGL Shading Language (GLSL), while DirectX 12 expects the High-Level Shader Language (HLSL) and a more structured offline compilation pipeline. Existing developer practice tends to fall into two categories: manual shader rewriting for tight control over hardware-specific optimizations, or cross-compilation workflows that preserve shader intent while reducing rewrite cost (Capodiecì, Cavicchioli, & Marongiu, 2022).

For a visually sensitive effect such as AnitoPlume’s volcanic plumes, the second approach is highly attractive. Manual translation risks introducing numerical drift and subtle visual discrepancies in lighting and blending logic. This design philosophy is exemplified by the broader shift in modern graphics tooling toward intermediate shader representations, which architecturally separate program descriptions from backend compilation. By utilizing intermediate representations like SPIR-V, such tooling prioritizes shader portability over hard-coded backend assumptions. In a migration study, this supports the argument that accurate visual porting is not just about translating syntax from GLSL to HLSL, but about maintaining shader behavior across compilation targets using tool-assisted pipelines like the DirectX Shader Compiler (DXC) or SPIRV-Cross (Kinkor, 2022).

2.8.2 Pipeline State Management Adaptation

A second major challenge is the shift from OpenGL’s dynamic state model to DirectX 12’s explicit Pipeline State Objects (PSOs). In architectures, OpenGL allows rendering states to be changed incrementally at runtime, with the graphics driver silently inferring and tracking state combinations (J. A. Shiraef, 2016). DirectX 12 completely eliminates this hidden driver-side management. Instead, explicit APIs require developers to bake the blend state, rasterizer state, depth-stencil state, primitive topology, and shader kernels into immutable PSOs ahead of time (Bossard, 2019).

The necessity of this structural rewrite is inherent to the explicit PSO model itself. Rather than relying on hidden driver-side state inference, explicit graphics APIs require applications to aggregate vertex layouts, shader kernels, rasterizer states, and framebuffer descriptions into a single, pre-compiled state object definition ahead of time. At draw time, the render context binds this monolithic object directly before issuing the call. For plume rendering, adopting this explicit design is critical; even minor mismatches in depth testing or alpha blending across different PSOs can fundamentally alter how overlapping particle billboards accumulate

and fade (Del Gallego et al., 2026). Therefore, PSO migration must be framed as a foundational architectural overhaul rather than a simple API substitution (Asplund, 2020).

2.8.3 Resource Synchronization and Memory Barriers

The most technically demanding aspect of the explicit migration process is resource synchronization. In OpenGL, hardware hazards are concealed by the driver, meaning developers rarely need to explicitly manage the transitions between writing to and reading from the same texture or buffer (Unterguggenberger et al., 2022). DirectX 12, by contrast, transfers the burden of memory management entirely to the developer, requiring explicit resource barriers to ensure that GPU accesses occur in the correct deterministic order (Rossant & Rougier, 2021).

This explicit synchronization burden is a defining characteristic of explicit graphics APIs, which require applications to issue resource barriers that route resources into their correct states before any updates, copies, or draw calls are executed. By manually handling transitions at both the resource and subresource levels, these architectures reinforce that state transitions are an application responsibility rather than an invisible driver service. For the AnitoPlume study, this supports a key argument: accurate migration is not just about producing a similar output image, but about guaranteeing that the GPU views each buffer in the correct state at the exact right microsecond. This is essential when updating thousands of particle billboards, as incorrect barrier placement during parallel draw calls leads to race conditions, stale texture reads, flickering artifacts, or catastrophic device crashes (Cheung, 2024).

Chapter 3

Computer Graphics and Rendering Concepts

The modernization of AnitoPlume involves migrating the rendering pipeline from a traditional to a modern graphics API architecture. This migration introduces a fundamentally different graphics architecture characterized by explicit resource management, reduced driver overhead, and improved utilization of modern multi-core processors. These architectural changes are beneficial for simulators such as AnitoPlume, which require continuous updates of thousands of dynamic particle billboards.

3.1 Graphics Rendering Architectures

Traditional graphics API architectures are designed to abstract low-level graphics operations from the programmer by implicitly managing graphics state, resource synchronization, and command validation within the graphics driver. Throughout this study, these are referred to as implicit graphics APIs, whereas architectures that expose these responsibilities directly to the application are referred to as explicit graphics APIs.

3.1.1 Overview of Graphics API Architectures

Implicit graphics APIs such as OpenGL are based on a state-based programming model where the application modifies global rendering states that are maintained

by the graphics implementation (“The OpenGL Graphics System: A Specification”, 2022). Implicit graphics API architecture relies on a thick driver layer that abstracts graphics pipeline management from the application (Fig. 9). Because the driver is responsible for maintaining graphics state, performing runtime validation, tracking resource usage, and managing command execution, significant CPU overhead is introduced during rendering (Ping et al., 2026).

Explicit graphics architectures expose synchronization, resource management, and command submission responsibilities directly to applications (Khronos Group, 2024). Rather than relying on implicit driver-managed state changes, explicit graphics architectures require the application to explicitly define and manage the configuration of the graphics pipeline. Applications explicitly define pipeline states, resource allocation, synchronization, and memory management as seen on figure 9.

OpenGL	Vulkan	DirectX 12
API originally created for graphics workstations; uses split memory.	Match for modern platforms (mobile/unified memory/tiled rendering).	Modern API for Windows/Xbox; scales across unified platforms.
Driver handles state validation/error tracking, limiting performance.	Explicit API gives application direct, predictable GPU control.	Shifting validation to driver, giving developers explicit resource control.
Obsolete single-threaded model; no parallel command generation.	Built for multi-core/threads; parallel command buffers.	Multi-core optimized; pre-computes command lists on any CPU core.
Complex choices; syntax evolved over 20+ years.	No legacy requirements; simpler API design and smaller specification.	Replaces legacy state objects with unified Pipeline State Objects (PSO).
In-driver GLSL compiler; requires shipping raw shader source.	SPIR-V target enables front-end language flexibility and reliability.	Compiles HLSL to intermediate DXIL/DXBC to reduce driver workload.
High vendor-specific implementation variability.	Simpler API and testing improve cross-vendor compatibility.	Strict Feature Levels (e.g., DX12 Ultimate) ensure vendor consistency.

Table 2: Graphics API comparison: OpenGL vs. Vulkan vs. DirectX 12 from J. Shiraef and Zoraja et al.

Table 2 summarizes the key architectural and performance differences between the OpenGL, Vulkan, and DirectX 12 graphics APIs. It contrasts the single-

threaded design of OpenGL with the modern, multi-threaded, and explicitly controlled paradigms introduced by Vulkan and DirectX 12.

The limitations of the traditional graphics API architecture become particularly significant in computationally intensive scientific visualization systems such as AnitoPlume. In addition to numerically integrating volcanic plume dynamics including gravity, buoyancy, drag, and wind forces for thousands of particles each simulation step, the CPU must also handle graphics driver validation and state management (Del Gallego et al., 2026). This added overhead limits rendering performance, motivating the use of lower-overhead graphics APIs to reduce CPU-side workload for applications that require high-throughput rendering. As rendering workloads become more complex, these driver-managed operations can limit the application’s ability to optimize resource usage, creating CPU-side bottlenecks that reduce the rate at which work can be submitted to the GPU (Ping et al., 2026).




Figure 9: Comparison of implicit graphics API architecture (left) with explicit graphics API architecture (right).

3.1.2 Performance Optimization opportunities

Explicit graphics architectures reduce driver-managed serialization by shifting the responsibility to the application. Tasks such as pipeline configuration, resource allocation, and command preparation can be performed explicitly by the application before or outside the main rendering loop. During runtime, the system only executes predefined instructions without modifying or accessing the underlying data. This shift allows command generation to be safely distributed across multiple CPU threads (Ping et al., 2026). The lightweight driver architecture introduced by explicit graphics APIs minimizes runtime validation and hidden state management, reducing CPU overhead while providing developers with greater control over rendering operations. By enabling application-driven optimization and more efficient utilization of hardware resources, explicit graphics API architectures are better suited for performance-critical rendering workloads (Ioannidis & Boutsis, 2020).

3.2 GPU Parallel Computing

3.2.1 General Purpose GPU Computing

The increasing computational demands of scientific visualization and physically based simulation have driven the adoption of Graphics Processing Units (GPUs) as general-purpose parallel computing devices. Unlike Central Processing Units (CPUs), which consist of a relatively small number of complex processing cores optimized for sequential execution, GPUs consist of many smaller cores organized into parallel execution units designed to maximize throughput for highly parallel workloads. As illustrated in Figure 10, this architecture enables GPUs to efficiently process workloads exhibiting high degrees of data parallelism, where identical computations are repeatedly applied across large collections of independent data elements (Owens et al., 2008).




Figure 10: CPU vs. GPU key differences. From (TecAdmin Staff, 2025).

The emergence of general purpose computing on GPUs transformed graphics hardware from fixed-function rendering accelerators into programmable parallel processors capable of executing scientific and engineering applications.

3.2.2 GPU Programming

Early GPU computing approaches utilized graphics APIs to express numerical algorithms as rendering operations. This paradigm evolved with the introduction of programming models such as NVIDIA’s CUDA, which exposed the GPU as a general parallel computing platform through familiar programming abstractions while providing direct access to the device’s computational resources (Luebke, 2008).

Harris describes GPU programming as a process of mapping computational problems onto massively parallel hardware by identifying operations that can be executed independently across many processing elements. Rather than accelerating a single complex computation, GPU applications achieve performance by decomposing a problem into numerous independent tasks that execute simultaneously. Consequently, algorithms exhibiting minimal inter-element dependencies and regular memory access patterns are generally well suited for GPU execution (Harris, 2004). Compute shaders organize execution into groups of parallel invocations that simultaneously execute the same program on different portions of the input data. Unlike vertex or fragment shaders, compute shaders are not

directly associated with rendering primitives, allowing applications to perform arbitrary numerical computations while explicitly managing memory resources and synchronization between execution stages (Khronos Group, 2024).

Because many scientific simulations consist of repeatedly updating large collections of independent entities, GPU computing has become widely adopted across computational physics, fluid dynamics, molecular simulation, image processing, and real-time visualization. Owens et al. identify data-parallel simulation as one of the primary application domains benefiting from GPU architectures due to their ability to execute thousands of lightweight computational threads concurrently. As GPU hardware and programming models continue to evolve, massively parallel computing has become a fundamental approach for accelerating computationally intensive scientific applications while maintaining interactive performance (Owens et al., 2008).

3.2.3 Application of Compute Shaders in Parallel Computing

Data-parallel simulation is one of the primary application domains benefiting from GPU architectures due to their ability to execute thousands of lightweight computational threads concurrently (Owens et al., 2008). Recent studies demonstrate the use of compute shaders across different simulation and visualization domains, including particle-based simulation, physical deformation models, and large-scale rendering.

Va, Choi, and Hong further demonstrated the use of compute shaders for iterative physics simulation, where simulation stability and performance depend on repeated updates across discrete time steps. Similarly, Madera-Ramirez, Lopez-Martinez, Moo-Mena, and Gomez-Montalvo applied compute shaders to particle-based sphere simulation by distributing particle updates and collision operations across GPU threads. Their work addressed the increasing computational and memory requirements caused by large particle counts by improving CPU–GPU buffer communication and maintaining a reusable computation buffer. Schütz, Kerbl, and Wimmer demonstrated the use of compute shaders for large-scale point cloud rendering. Their approach replaced portions of the traditional graphics pipeline with compute-based processing, allowing independent point data to be processed through GPU parallel execution. Although these works address different applications, they illustrate how compute shaders exploit data parallelism to improve throughput for workloads consisting of many independent elements.

Table 3: Application of compute shaders for optimization.

Study	Domain	Compute Shader Application	Results
Va et al. (2021)	Real-time cloth simulation	Implemented a GPU-based parallel cloth simulation using compute shader kernels for mass-spring calculations and constraint enforcement	GPU parallel implementation achieved faster performance than CPU-based implementation while maintaining stable cloth behavior using constraint-based methods.
Madera-Ramírez et al. (2021)	Particle Simulation	Used compute shaders to parallelize particle updates, collision operations, and GPU buffer communication	Achieved up to 100× speedup over the sequential implementation and enabled processing of larger particle counts.
Schütz et al. (2021)	Point cloud rendering	Used compute shaders to perform point processing and rendering operations through GPU-based parallel execution	Demonstrated high-throughput rendering of hundreds of millions of points with improved performance over previous approaches.

Similar characteristics are present in particle-based plume simulation, where each plume element maintains independent state variables that must be repeatedly updated during each simulation step. Therefore, compute shaders provide a suitable computational model for accelerating both simulation updates and visualization of large numbers of dynamic plume elements.

3.3 Lagrangian and Eulerian Models for Volumetric Plume Dynamics

3.3.1 Lagrangian Formulation

AnitoPlume utilizes a Lagrangian formulation for representing volcanic plume evolution, where the plume is modeled as a collection of discrete volumetric elements whose physical states are updated over time. Unlike Eulerian approaches that evaluate flow properties within fixed spatial domains, Lagrangian approaches track the movement and evolution of individual elements within the flow field (Esposti Ongaro et al., 2020).

The plume formulation presented by Lastic et al. (2022) and implemented in AnitoPlume (Del Gallego et al., 2026) represents plume dynamics through discrete elements with associated physical properties. Each plume element maintains state variables such as position, velocity, mass, density, and volume. During each simulation step, these variables are updated based on governing physical equations that consider forces including gravity, buoyancy, atmospheric drag, and wind effects (Lastic et al., 2022; Del Gallego et al., 2026; Folch, Costa, & Macedonio, 2016).

The structure of the Lagrangian formulation exhibits data-parallel characteristics because individual plume elements can be updated primarily using their local state information and shared simulation parameters. Consequently, operations such as force evaluation and numerical integration can be distributed across multiple processing units, provided that synchronization requirements are properly managed (Lastic et al., 2022).

However, Lagrangian particle-based methods also present limitations when applied to large-scale geophysical phenomena. The computational cost associated with representing the large number of particles required for realistic volcanic processes has limited the widespread use of fully Lagrangian approaches in volcanological simulations (Esposti Ongaro et al., 2020). These limitations motivate the consideration of alternative formulations, model refinements, or modifications to the underlying mathematical representation depending on the accuracy and computational requirements of the simulation. Consequently, the selection of governing equations and the formulation of the plume model remain important considerations in the continued development of volcanic simulation systems.

Despite these considerations, the data-parallel characteristics of the Lagrangian representation make plume simulation suitable for optimization through modern

parallel processing architectures. Since the computational workload consists of continuously updating large numbers of dynamic elements, improving execution efficiency requires an architecture capable of reducing overhead and efficiently utilizing available hardware resources. Explicit graphics architectures provide mechanisms that support this objective by enabling greater control over resource management and parallel command execution.

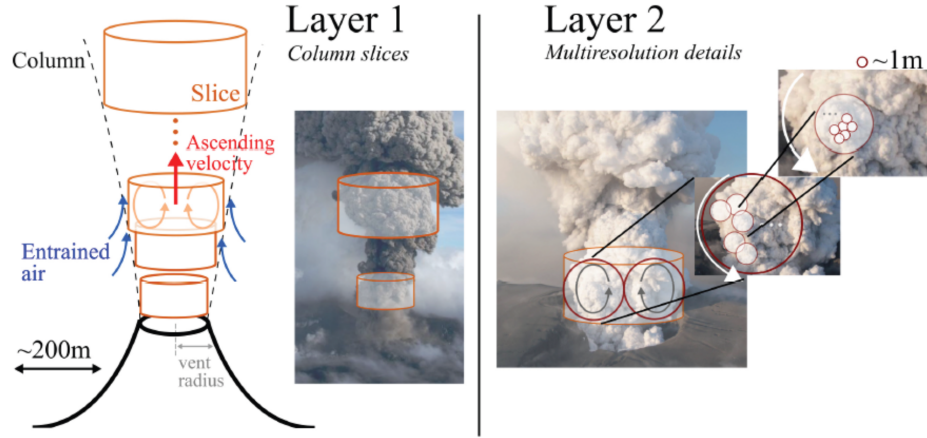


Figure 11: Lagrangian approach. From (Lastic et al., 2022).

Figure 11 illustrates the Lagrangian formulation implemented by (Lastic et al., 2022), detailing the physical forces, wind-driven trajectory changes, and volumetric expansion acting on individual discrete plume elements over time.

3.3.2 Eulerian Formulation

In contrast to the Lagrangian approach, the Eulerian formulation describes plume dynamics with reference to fixed points or cells in space, rather than tracking discrete elements as they move through the flow field. Under this formulation, the volcanic plume is represented as a continuous field of quantities, such as ash concentration, density, temperature, and velocity, which are computed and updated at fixed grid locations over time as material flows through them (Esposti Ongaro et al., 2020). Rather than following individual particles or plume elements as they travel through the atmosphere, the Eulerian approach numerically solves governing transport equations directly on a spatial grid that discretizes the simulation domain. This formulation underlies many of the numerical geoscientific models used for volcanic hazard forecasting; for instance, models such as FALL3D use

Eulerian atmospheric dispersion techniques to compute the transport, dispersion, and eventual deposition of volcanic tephra across a defined spatial domain (Folch et al., 2009; Del Gallego et al., 2026). While this grid-based approach is generally well-suited for capturing diffusive processes and producing quantitatively accurate forecasts, it comes at a significant computational cost, particularly as grid resolution increases or the simulation domain grows, since the governing equations must be solved at every grid cell regardless of whether volcanic material is present at that location. Consequently, Eulerian volcanic transport models are typically designed for offline scientific analysis and hazard mapping rather than for real-time or interactive visualization (Del Gallego et al., 2026).




Figure 12: Eulerian approach. From (Folch et al., 2020).

For the Eulerian Approach, Figure 12 visualizes how it discretizes the simulation space into a 3D staggered mesh of fixed computational cells to evaluate mass transport and flow velocity over time.

3.3.3 Comparison of Both Approaches

AnitoPlume adopts the Lagrangian formulation because it prioritizes real-time interactivity over the grid-based accuracy of Eulerian models. This study preserves that formulation as a fixed theoretical foundation, since its scope is limited to the rendering and software architecture layer, not the underlying plume dynamics. The Lagrangian model’s data-parallel characteristics, individual elements updating largely independently using local state, are what make it a suitable target for

Table 4: Summary of key theoretical distinctions between the Lagrangian and Eulerian formulations.

Dimension	Lagrangian Formulation	Eulerian Formulation
Frame of reference	Tracks discrete, moving plume elements.	Evaluates fixed spatial grid points.
Representation	Discrete particles with individual states (pos, vel, mass).	Continuous field of quantities (concentration, density).
Computation	Scales with particle count; local-state updates.	Scales with grid size; solved at every cell.
Accuracy	Approximates diffusive behavior.	High quantitative accuracy for diffusive transport.
Real-time	Well-suited; lightweight for moderate counts.	Generally unsuitable; computationally demanding.
Representative use	Lastic et al. (2022); Anito-Plume	FALL3D (Folch et al., 2009)
Limitation	Cost grows with particle count for visual realism.	Cost grows with grid resolution/domain coverage.

the proposed low-level rendering architecture, which is designed to expose fine-grained control over parallel command execution and GPU resource management. This theoretical basis justifies why the present study’s contributions are directed at the rendering pipeline rather than the plume simulation model itself.

3.4 Profiling Tools

3.4.1 The Role of Profiling in Graphics Benchmarking

Profiling is the process of instrumenting a software system to collect quantitative data about its runtime behavior: execution time, resource utilization, throughput, and synchronization patterns (Cheung, 2024). The purpose is to identify performance bottlenecks and validate optimization efforts. In graphics development, profiling operates across two domains: the CPU side, where the application records and submits commands, and the GPU side, where those commands are translated into rendering and compute workloads (Ping et al., 2026). For scientific visualization applications such as AnitoPlume, where a single frame may involve thousands of particle billboard draws, a terrain pass, and a UI composite, profiling is the only means of determining precisely where frame time is being consumed.

In implicit graphics APIs, profiling is primarily concerned with identifying

draw-call overhead generated by heavy driver validation and state tracking (J. Shiraef, 2016). In explicit graphics APIs, the profiler surfaces a broader set of implementation concerns: the cost of recording command buffers, the impact of pipeline barrier placement on GPU occupancy, the efficiency of descriptor set layouts, whether pipeline state objects are reused effectively, and whether compute-to-render transitions are properly scheduled to avoid stalling the pipeline (Ping et al., 2026; Khronos Group, 2024). These concerns go beyond the quality of the rendered output, and failure to profile them results in a system that may produce correct visuals but fall short of the hardware’s capabilities (Ioannidis & Boutsis, 2020).

The following subsections survey three profiling tools relevant to explicit API development and benchmarking. These tools differ in scope, ranging from real-time statistical profiling to offline frame debugging to deep hardware performance analysis, and together illustrate the instrumentation dimensions that modern graphics development demands.

3.4.2 Tracy Profiler

Tracy is an open-source, cross-platform profiler that works across all GPU hardware vendors and is designed for real-time, frame-based performance analysis (Bartosz Taudul, n.d.). It requires manual source-code instrumentation: the developer inserts zone-marking macros around blocks of interest, such as the simulation update loop, command buffer recording, and presentation. Tracy records the start time, duration, and thread affinity for each zone across every captured frame, building statistical profiles over time.

On the GPU side, Tracy supports recording timestamp queries directly within command buffers, allowing the profiler to track GPU-side execution activity on the same timeline as CPU zones (Bartosz Taudul, n.d.). This integration is significant for evaluating parallelization: the developer can observe whether a worker thread finishes updating its assigned smoke layers well before others, indicating load imbalance, or whether GPU execution overlaps with CPU command recording rather than idling while waiting for CPU submission.

Tracy is suited for identifying trends that manifest over multiple frames, such as frame timing spikes, CPU-side bottlenecks in the physics integration pass, and idle gaps caused by thread barriers. It is less suited for per-event GPU-level diagnostics, which require offline inspection of individual frames.

3.4.3 NVIDIA Nsight Graphics

NVIDIA Nsight Graphics is a GPU profiling and debugging tool that provides access to low-level hardware and performance counters available on NVIDIA GPUs (NVIDIA Corporation, 2023). Unlike generic frame capture tools, Nsight Graphics directly queries hardware performance counters including shader processor occupancy, texture unit utilization, memory bandwidth throughput, L1 and L2 cache hit rates, and warp stalls for each draw or dispatch call within a captured frame.

These metrics provide resolution that a generic frame-capture profiler cannot offer. For example, identifying that a fragment shaded render pass for volcanic plume billboards is bound by texture memory latency rather than arithmetic throughput suggests specific optimization steps: reducing texture resolution, packing multiple data channels into a single texture, or altering the filtering mode. Each step can be validated against the same hardware counter data.

Because Nsight Graphics is restricted to NVIDIA hardware, it is most appropriate for detailed single-frame profiling on NVIDIA workstations rather than continuous monitoring across diverse hardware. For research groups that deploy NVIDIA GPU workstations, Nsight provides the deepest constituent view of hardware performance available.

3.4.4 RenderDoc

RenderDoc is a free, open-source graphics debugging tool that supports Vulkan, DirectX 12, DirectX 11, and OpenGL across GPUs from NVIDIA, AMD, and Intel, making it the most broadly applicable tool surveyed here (Baldur Karlsson, 2023). It captures the full sequence of graphics API calls for a single user-specified frame, serializing every draw call, compute dispatch, render pass transition, and memory barrier into an offline trace.

The capture workflow is straightforward: the application is launched through RenderDoc’s process injector, a frame is captured via the overlay or hotkey, and the resulting trace is opened in the inspection interface. That interface provides the following:

- **Event Browser.** The time-ordered list of every API call in the frame, allowing the developer to verify that command sequences, barriers, and render pass transitions are executed in the intended order.
- **Pipeline State View.** The complete GPU state, including bound shaders,

descriptor sets, vertex and index buffers, rasterizer state, depth-stencil state, and blend state, for any selected event. This enables verification that Pipeline State Objects are correctly constructed and bound.

- **Resource Inspector.** A per-pixel visualization of any intermediate render target, texture, or buffer at any point inside the frame. This is essential for confirming that the terrain pass, plume particle pass, and UI composite pass each produce correct intermediate outputs.
- **GPU Timing.** Per-event GPU-side execution time derived from timestamp queries, identifying which draw calls or compute dispatches are the most expensive within a single frame.

3.4.5 Summary of Profiling Tools

Each of the profiling tools discussed addresses a distinct layer of graphics performance analysis, and their combined use provides comprehensive instrumentation coverage for projects migrating between graphics APIs. Table 5 below contrasts them across five dimensions particularly relevant to this study.

Table 5: Comparative summary of the profiling tools discussed.

Dimension	Tracy	NVIDIA Nsight Graphics	RenderDoc
Profiling Mode	Real-time, continuous across frames	Single-frame, offline	Single-frame, offline
Coverage	CPU zones and GPU timer queries on shared timeline	Per-event hardware counter metrics (occupancy, cache, bandwidth)	Full API call trace, pipeline state, GPU timing
GPU Vendor Support	All (AMD, Intel, NVIDIA)	NVIDIA only	All (AMD, Intel, NVIDIA)
API Support	Any (via vendor-ext timestamp queries)	Vulkan, DirectX 12, DirectX 11, OpenGL	Vulkan, DirectX 12, DirectX 11, OpenGL
Primary Use Case	Identifying CPU-side bottlenecks over multiple frames; validating multi-threaded parallelism	Diagnosing hardware-level stalls (cache misses, warp divergence, low occupancy)	Inspecting pipeline state correctness and per-event GPU timing across APIs

Chapter 4

Methodology

The primary objective of this research is to port and extend AnitoPlume’s real-time simulation framework by modernizing its rendering pipeline, optimizing hardware utilization through new, explicit graphics APIs, and refining its underlying physical simulation models.

4.1 Implementation and Migration

4.1.1 Legacy Simulation

The legacy execution architecture relies on a strictly sequential, single-threaded pipeline driven by an iterative time-stepping loop. Within each frame slice, the simulation runs a high-frequency sub-stepping loop (executing ten physics updates per render frame) to remove visual smoke jumping and oscillations due to numerical instability. The control flow inside this sub-step operates linearly across distinct global entity arrays.

As shown in Figure 13, the simulation is organized around the smoke layer as the primary simulation unit. Each smoke layer models the behavior of the eruption column by computing atmospheric air entrainment, thermodynamic properties (e.g., temperature and density), external forces such as buoyancy, gravity, wind, and drag, and integrating the resulting motion over time. Attached to every smoke layer are free spheres, which represent smoke that inherit the motion and physical state of their parent layer while evaluating transitions into falling or neutrally buoyant states. Each free sphere further generates a hierarchy of subspheres that

provide the volumetric detail required for visualization. These subspheres are rendered as billboard sprites to produce the final volumetric appearance of the smoke cloud but do not participate in the physical simulation.




Figure 13: Legacy physics simulation loop.

AnitoPlume initializes the simulation environment by loading a preprocessed three-dimensional terrain model representing Taal. The terrain mesh is generated offline from geospatial elevation datasets and optimized for interactive rendering while preserving the topographic features required for volcanic hazard visualization (Del Gallego et al., 2026). Textures for the terrain model were made through photobashing of satellite images taken from Google Earth; as seen on Figure 14 (Del Gallego et al., 2026).




Figure 14: Workflow for texturing the surrounding area of Taal Volcano. From (Del Gallego et al., 2026)

During initialization, the terrain is imported into the rendering pipeline, transformed into the simulation coordinate system, and associated with its corresponding surface texture.




Figure 15: Legacy terrain rendering pipeline.

In addition to visualization, the terrain serves as the physical boundary of the simulation. Height and surface normal information are queried during particle

updates to detect terrain intersections and compute particle-ground interactions, allowing falling volcanic material to respond to the underlying topography. Consequently, the terrain functions as both the rendered environment and the collision surface used by the simulation.

Beyond the core simulation and terrain rendering pipelines, AnitoPlume also provides an interactive graphical user interface (GUI) through which users can manipulate plume behavior in real time. This interface additionally surfaces contextual 3D tooltips that display relevant information about Taal Volcano directly within the scene. Figure 16 shows this interface, representing the primary means by which users interact with and observe the legacy simulation system.




Figure 16: GUI of AnitoPlume for manipulating plume behavior and 3D UI tooltips for displaying Taal Volcano information to the user. From (Del Gallego et al., 2026)

4.1.2 Parallelization of the Simulation Loop

One possible improvement for AnitoPlume’s performance is to explore multi-threaded rendering. To achieve this, we propose implementing parallelization to improve the responsiveness of the interactive simulation for multicore CPUs. This concurrent execution is feasible because individual smoke layers operate in complete isolation from one another. The mathematical functions executed during the smoke layer update phase do not feature any cross-layer spatial dependencies; instead, the physics equations for any given layer depend strictly on a shared set of initial environmental parameters, vent parameters, and global trackers. Because

later simulation stages consume data produced by earlier stages, we propose introducing barrier synchronization between execution phases, as shown in Figure 17. These synchronization points will guarantee that all worker threads complete one phase before subsequent phases begin, ensuring that rendering is performed only after all simulation stages have concluded to maintain visual consistency.




Figure 17: Parallelization of smoke layer update phase.

Beyond CPU multithreading, the same data-parallel characteristics also make the simulation suitable for execution on modern GPUs using compute shaders. Hence, a further performance optimization for AnitoPlume is to offload the simulation loop to the GPU via dedicated compute workloads.

Adopting the GPU-accelerated design pattern proposed by Madera-Ramirez et al., the simulation and rendering tasks are decoupled into distinct compute and render stages. As illustrated in Figure 18, the CPU is relieved of per-layer physics calculations and serves primarily as an orchestrator which loads global parameters and dispatches GPU workloads. Simulation data is stored in GPU resident memory buffers where compute passes execute the numerical update functions for each smoke layer concurrently. Explicit synchronization barriers will enforce that all GPU compute passes complete before the update buffers are read by the rendering pipeline.




Figure 18: Proposed parallel architecture.

By exploiting both multicore CPUs and GPU parallelism where applicable, the proposed execution architecture aims to reduce the computational bottlenecks of the legacy sequential implementation while preserving the underlying Lagrangian plume model.

4.1.3 Terrain Rendering Improvements

While the terrain geometry will accurately capture the topography of the study area, the initial rendering implementation will be intentionally simple. We will render the terrain using a single textured mesh with basic diffuse lighting and fixed material properties, without terrain-specific rendering techniques such as multi-layer texture blending, physically based shading, normal mapping, ambient occlusion, or level-of-detail management.

The proposed rendering architecture will enable more advanced rendering techniques to be incorporated without modifying the underlying physics simulation, since we will manage rendering resources independently through explicit GPU resource management. One promising approach is to utilize this architecture to facilitate the integration of enhancements such as higher-resolution terrain textures, physically based materials, dynamic lighting, normal mapping, shadow mapping, terrain tessellation, and level-of-detail rendering while preserving the existing sim-

ulation workflow.

Although we will modernize the terrain rendering pipeline, the existing graphical user interface as described in Figure 16 will remain intact. We will preserve user interactions for terrain visualization, simulation control, and parameter configuration, ensuring compatibility with the original application workflow.

4.2 Benchmarking Procedure

4.2.1 Profiling Tool Selection

As discussed in Section 3.4, profiling tools for explicit graphics API development fall into several categories: real-time CPU-GPU profilers, hardware-level performance analyzers, and frame capture debuggers. Each offers a different scope of instrumentation. The selection of a primary profiling instrument for this study requires that the tool satisfy three conditions simultaneously:

1. **Cross-vendor hardware support.** The benchmarking procedure is designed to produce reproducible results across a range of consumer and workstation GPUs from NVIDIA, AMD, and Intel. A profiling instrument locked to a single GPU vendor would limit the generalizability of the results and introduce hardware-specific reporting biases.
2. **Framework-agnostic comparison.** Because this study evaluates both the legacy OpenGL implementation in Phase 1 and the ported explicit API pipeline in Phase 3, the benchmarking methodology must remain independent of any specific rendering framework. Both pipelines will be profiled under identical hardware and workload conditions, and the same profiling tool must capture both runs using the same metric definitions and output formats. This approach ensures that performance comparisons reflect architectural differences rather than artifact of tooling or framework-specific optimizations.
3. **Evaluation metrics.** The tool must expose per-event GPU execution time, shader-stage latency, and memory bandwidth usage to populate the metrics defined in Table 6.

RenderDoc (Baldur Karlsson, 2023) is selected as the primary profiling instrument for this study. While Tracy (Section 3.4.2) identifies multi-frame trends

such as thread synchronization and CPU-GPU overlap, it records only aggregated zone-level timing and cannot isolate per-event draw calls or inspect individual pipeline states. RenderDoc’s offline frame-capture model provides the per-event granularity that Tracy lacks which is necessary to populate the metrics in Table 6, particularly Shader Execution Time, Command Recording Time, and GPU Memory Bandwidth. RenderDoc also supports the API used by the original AnitoPlume (OpenGL) and modern graphics APIs such as but not limited to Vulkan and DirectX 12 under an identical capture workflow, ensuring that profiles from Phases 1 and 3 are methodologically compatible and directly comparable, and it operates across NVIDIA, AMD, and Intel GPUs, unlike NVIDIA Nsight Graphics (Section 3.4.3) which is restricted to a single vendor and would limit the generalizability of benchmarking results.

4.2.2 Evaluation Metrics

To evaluate scalability, performance metrics will be collected through increasing simulation workloads, using profiling tools such as RenderDoc to capture low-level pipeline data alongside application-level timing instrumentation. The specific metrics utilized for this evaluation, along with their respective descriptions, are detailed in Table 6.

Table 6: Performance metrics and descriptions used for evaluation.

Metric	Description
Average FPS	Overall rendering performance
Frame Time (ms)	Total frame execution time
CPU Frame Time	CPU workload
GPU Frame Time	GPU workload
Command Recording Time	CPU time spent preparing GPU commands
Shader Execution Time	GPU shader workload
GPU Memory Bandwidth	Memory utilization
CPU Utilization	Effectiveness of multithreading

Workload size and rendering performance will be visualized using line graphs illustrating trends in average frame rate, frame time, CPU utilization, and GPU utilization. GPU profiling metrics, including shader execution time, command recording time, and memory bandwidth utilization, will likewise be compared to determine how efficiently the proposed architecture utilizes available hardware resources.

Finally, the benchmarking results will be interpreted alongside the User Acceptance Testing (UAT) results to determine whether improvements in computational

performance correspond to improvements in perceived responsiveness and overall user experience.

4.3 Usability Testing

To verify that the architectural overhaul maintains or improves the interactive user experience, an exploratory usability test will be conducted with target beneficiaries:

- **Participants:** Recruitment will follow a tiered priority order. PHIVOLCS personnel will be the top priority given their domain expertise in volcanology and direct familiarity with Taal Volcano hazard assessment. Should this pool be insufficient, recruitment will expand to, in order: (1) individuals who have witnessed a Taal Volcano eruption firsthand, (2) enthusiasts with experience using volcano/geohazard simulation software, and (3) computer science students with a background in rendering or volcano-related coursework/projects. This staged recruitment reflects the tool's intended end-users while ensuring an adequate sample size, and informs the corresponding participant categories documented in the study's ethics clearance and consent materials.
- **Setup:** After screening and obtaining consent from the Informed Consent Form (ICF), participants will be given time to explore the interface.
- **Task:** Participants will use the software to recreate specific Taal Volcano eruption plumes, following a set of predefined scenarios.
- **Survey:** Users will complete the System Usability Scale (SUS) questionnaire to rate the software's ease of use.
- **Feedback:** A closing interview will gather suggestions for improving the interface and overall performance.

Table 7: Sample SUS questionnaire implemented by (Del Gallego et al., 2026). 1 — strongly disagree, 5 — strongly agree.

Code	Question
Q1	I think that I would use AnitoPlume frequently.
Q2	I found AnitoPlume unnecessarily complex.
Q3	I found AnitoPlume easy to use.
Q4	I think that I would need the support of a technical person to be able to effectively and efficiently use AnitoPlume.
Q5	I found that various features of the simulator were well integrated in AnitoPlume.
Q6	I thought there was too much inconsistency in AnitoPlume.
Q7	I would imagine that most people without prior experience in using similar software would easily learn how to use AnitoPlume.
Q8	I found AnitoPlume very cumbersome (slow or complicated and therefore inefficient) to use.
Q9	I felt very confident while using AnitoPlume.
Q10	I needed to figure out a lot of things (such as hotkeys, navigating the UI) before I could confidently use AnitoPlume.

4.4 Research Workflow

This study follows a three-phase research workflow, outlined in Table 8, that guides the porting and extension of AnitoPlume from its legacy rendering implementation to the proposed rendering architecture, from initial legacy analysis through final evaluation of performance and usability.

Phase	Research Activity	Purpose / Output
1	Legacy Pipeline Profiling and Bottleneck Analysis	Evaluate the existing simulation and rendering architecture to establish baseline performance metrics. Profile CPU-GPU synchronization, draw call overhead, and frame pacing to identify compute and graphics pipeline bottlenecks that impede scalability.
2	Architectural Refactoring and Pipeline Optimization	Design and iteratively implement optimizations within the execution model and rendering pipeline. This involves exploring modern graphics paradigms (e.g., explicit command submission, reducing driver overhead) and evaluating concurrent execution strategies to maximize render throughput, adapting the architecture based on empirical performance gains.
3	Performance Benchmarking and Interactive Evaluation	Conduct a comparative analysis between the legacy and refactored pipelines using low-level GPU profiling metrics (e.g., shader execution time, memory bandwidth utilization). Perform user testing to correlate these hardware-level optimizations with perceived interactive responsiveness and visual consistency.

Table 8: Overview of the research phases, activities, and objectives.

Appendix A

Research Ethics Documents

A.1 Clearance Form



RESEARCH / PROJECT ETHICS REVIEW FORM For Thesis / Capstone / Dissertation Proposals (ver.Oct2024c)		
Names of Student Researcher(s): Adrian Rafael Bernandino Anton Miguel Borromeo Russell Emmanuel Galan Reinman Geoffrey Ganayo		
Research/Thesis/Capstone Title: A New Rendering and Plume Simulation Framework for Visualizing Taal Volcano Eruption Plumes		
College: College of Computer Studies		
Department: Software Technology		
Course: BS Computer Science major in Software Technology		
Expected Duration of the Project: from: Aug 6, 2026 to Aug 6, 2027		
PROTOCOL: <i>(To be filled up by proponents; Use the Ethics Checklists as guides in determining areas for ethical concern/consideration. Attach and properly reference the accomplished ethics checklists and supporting documents.)</i>		
		Reference Document (include page number)
Target participants	Sample Size: 10 to 20 voluntary participants, recruited according to a tiered priority order: PHIVOLCS personnel as the primary target group, followed by individuals who witnessed a Taal Volcano eruption firsthand, then volcano/geohazard simulation enthusiasts, and lastly computer science students with a background in rendering or volcano-related coursework, should earlier priority groups be insufficient to reach the target sample size. Inclusion Criteria: Recruitment will follow a tiered priority order. Priority 1: PHIVOLCS personnel with domain expertise in volcanology and hazard assessment. Priority 2 (if Priority 1 pool is insufficient): individuals who witnessed a Taal Volcano eruption firsthand. Priority 3: enthusiasts with experience using volcano/geohazard simulation software. Priority 4: computer science students or developers with a relevant	Scope and Limitations (page 4-5)

	background in 3D graphics, rendering, or volcano-related coursework/projects. Exclusion Criteria: Individuals who do not fall into any of the four priority categories above, or who are unable to complete the informed consent process. Withdrawal Criteria: Participants may exit the software at any stage and can choose to withdraw from the evaluation without academic or personal penalty.	
Data to be used / collected	Technical Data: Empirical hardware performance logs including Frame Rate (FPS), CPU/GPU resource utilization percentage, and frame rendering times (ms). Human Data: Quantitative usability scores via the System Usability Scale (SUS) questionnaire and qualitative feedback regarding visual artifacts or rendering anomalies. No personally identifiable information (PII) such as names, student numbers, or email addresses will be tied to the responses. Also, feedback gathered from a closing interview regarding suggestions for improving the interface and overall performance.	Scope and Limitations (page 4-5)
Brief procedure	Screening: Interested participants complete a short pre-screening form confirming which priority category they fall under (PHIVOLCS personnel, Taal eruption witness, volcano/geohazard simulation enthusiast, or CS student/developer with relevant graphics background), consistent with the tiered recruitment order. Consent: Eligible participants are presented with the Informed Consent Form (ICF), which outlines the evaluation's scope, hardware interactions, session length, and data privacy measures. Orientation/Briefing: Participants are given a brief, neutral walkthrough of the	Scope and Limitations (page 4-5)

	interactive Taal Volcano Plume simulator's controls and features. 4. Usability Task: Participants will use the software to recreate specific Taal Volcano eruption plumes in response to predefined scenarios. 5. Post-Session Questionnaire : Participants will complete an anonymous, self-administered SUS questionnaire to rate the software's ease of use and compare visual consistency. 6. Feedback: A closing interview will be conducted to gather suggestions for improving the interface and overall performance.	
Potential risks	Physical: Minimal risks, though participants may experience minor eye strain or visual fatigue from interacting with the 3D graphics simulator. Technical: Participants may experience frustration if they encounter bugs or usability issues with the experimental prototype. Mitigation: Researchers will provide immediate technical support for major difficulties, and participants are encouraged to request breaks at any time if they feel tired or frustrated. Psychological: For participants recruited as eruption witnesses, engaging with a simulation recreating the Taal eruption may evoke discomfort or distress related to their lived experience. Mitigation: The ICF for this participant category will explicitly disclose the nature of the simulated content in advance, participants will be reminded they may stop or withdraw at any point without penalty, and researchers will monitor for visible signs of distress during sessions.	

Applicable Terms of Use or Licensing Agreement of Data, Models and/or Tools	AnitoPlume Source Code: Used with explicit permission and under the supervision of the thesis adviser, Neil Del Gallego. Libraries: Eigen (Mozilla Public License v2.0) and dearIMGUI (MIT License). Graphics Toolkits & SDKs: Khronos Vulkan SDK, Microsoft DirectX 12 SDK, and Khronos OpenGL specifications. Profiling Tools: RenderDoc (MIT License).	Scope and Limitations (page 4-5)
Steps to safeguard the participants and/or data	Participant Safety: Testing sessions will be kept brief to prevent eye strain. Participants will be explicitly informed via the ICF that participation is entirely voluntary and that they may withdraw from the usability test at any point without academic or personal penalty. Participants recruited as PHIVOLCS personnel will be engaged through formal coordination with the agency, with a permission letter obtained and session scheduling arranged to minimize disruption to their official duties. Data Safety: All performance logs and usability questionnaire responses will be fully anonymized using alphanumeric identifiers. Data Storage: Survey sheets and digital spreadsheets will be secured in a password-protected cloud storage repository accessible only to the primary student researchers. Data Retention: Anonymized data will be kept 4 to 5 years following final publication or degree conferral to allow for reproducibility audits before permanent deletion	Scope and Limitations (page 4-5)
Other concerns and considerations	None at this point in time	
PANEL'S ASSESSMENT/COMMENTS:		

PANEL'S DECISION:

- € **Approved** (*Proponent/s may proceed with the study/data collection*)
- € **Modifications Required** (*Proponent/s should address ethical issues, follow the panel's recommendations for required modifications, and resubmit Ethics Review Form for approval before proceeding with the study or before starting with data collection*)
- € **For Full Review** (*Proponent/s should submit the proposal to the University Research Ethics Review Committee (RERC) and obtain a research ethics clearance before proceeding with the study or before starting with data collection*)

JUSTIFICATION:

Name and Signature of Adviser/Mentor

Date: _____

Name and Signature of Panel Member

Name and Signature of Panel Member

Date: _____

Date: _____

_____ _____ _____ Name and Signature of Panel Member Date: _____
--

A.2 General Checklist

 De La Salle University	Research Ethics Review Committee Research Ethics Office, 3F Henry Sy Sr. Hall De La Salle University Manila 2401 Taft Avenue, Manila 1004, Philippines REO@dlsu.edu.ph (632) 524-4611 loc. 513	SOP No.: 2
		Form No.: 2(D)
		Version No.: 1
		Version Date: July 2016

DE LA SALLE UNIVERSITY General Research Ethics Checklist
<i>This checklist is to ensure that the research conducted by the faculty members and students of De La Salle University is carried out according to the guiding principles outlined in the Code of Research Ethics of the University. The investigator is advised to refer to the <u>De La Salle University Code of Research Ethics and Guide to Responsible Conduct of Research</u> before completing this checklist. Statements pertinent to ethical issues in research should be addressed below. The checklist will help the researcher/s and advisers/readers/evaluators determine whether procedures should be undertaken during the course of the research to maintain ethical standards. The University's <u>Guide to the Responsible Conduct of Research</u> provides details on these appropriate procedures.</i>

Researcher Details	
Students	BERNANDINO , Adrian Rafael DC. BORROMEO , Anton Miguel G. GALAN , Russell Emmanuel G. GANAYO , Reinman Geoffrey A.
Thesis Adviser	DEL GALLEG0 , Neil
Department/College	Department of Software Technology / College of Computer Studies
Proposed Title of the Research	A New Rendering and Plume Simulation Framework for Visualizing Taal Volcano Eruption Plumes
Term(s) and academic year in which research project is to be undertaken	Term 3 of AY 25-26 to Term 3 of AY 26-27

<i>This checklist must be completed AFTER the De La Salle University Code of Ethics has been read and BEFORE gathering data.</i>		
Questions	Yes	No
1. Does your research involve human participants (this includes new data gathered or using pre-existing data)? If your answer is yes , please answer Checklist A (Human Participants) . Please specify if the kind of research you will be conducting falls under any of the following Human Participants sub-categories:	✓	
1.A. Will you be conducting Action Research in an existing business, company, or school? If your answer is yes , please answer Checklist F (Action Research) .		✓

 De La Salle University	Research Ethics Review Committee Research Ethics Office, 3F Henry Sy Sr. Hall De La Salle University Manila 2401 Taft Avenue, Manila 1004, Philippines REO@dlsu.edu.ph (632) 524-4611 loc. 513	SOP No.: 2
		Form No.: 2(D)
		Version No.: 1
		Version Date: July 2016

1.B. Does your research involve online communities (this includes culling data from social media platforms, online forums and blogs)? If your answer is yes , please answer Checklist G (Internet Research) .		✓
1.C. Does your research involve human participants who are situated in a community and may necessitate permission to acquire access to them? If your answer is yes , please answer Checklist H (Community Research) .		✓
2. Will your research make use of documents which are not in the public domain and, thus, require permission for use from the custodian of such documents? If YES, please provide certification that permission from the custodian of the data was sought and granted.		✓
3. Will your research make use of secondary data (e.g., surveys, inventories, plans, official documents, etc.) from an institution, organization, or agency, which are not in the public domain and, thus, require permission for use from the custodian of such documents? If YES, please provide certification that permission to use the data was sought from the institution, organization, or agency and approval was granted.		✓
4. Does your research involve animals (non-human subjects)? If your answer is yes , please answer Checklist B (Animal Subjects) .		✓
5. Does your research involve Wildlife? If your answer is yes , please answer Checklist C (Wildlife) .		✓
6. Does your research involve microorganisms that are infectious, disease causing or harmful to health? If your answer is yes , please answer Checklist D (Infectious Agents) .		✓
7. Does your research involve toxic/chemicals/ substances/materials? If your answer is yes , please answer Checklist E (Toxic Agents) .		✓

Research with Ethical Issues to address:

If you have a YES answer to any of the above categories, you will be required to complete a detailed checklist for that particular category. A YES answer does not mean the disapproval of your research proposal. By providing you with a more detailed checklist, we ensure that the ethical concerns are identified so these can be addressed in adherence to the University Code of Ethics.

 De La Salle University	Research Ethics Review Committee Research Ethics Office, 3F Henry Sy Sr. Hall De La Salle University Manila 2401 Taft Avenue, Manila 1004, Philippines REO@dlsu.edu.ph (632) 524-4611 loc. 513	SOP No.: 2
		Form No.: 2(D)
		Version No.: 1
		Version Date: July 2016

Declaration of Conflict of Interest

☒ 1. I do not have a conflict of interest in any form (personal, financial, proprietary, or professional) with the sponsor/grant-giving organization, the study, the co-investigators/personnel, or the site.

☐ 2. I do have a conflict of interest, specifically:

☐ A. I have a personal/family or professional interest in the results of the study (family members who are co-proponents or personnel in the study, membership in relevant professional associations/organizations).

Please describe the personal/family or professional interest:

☐ B. I have proprietary interest vested in this proposal (with the intent to apply for a patent, trademark, copyright, or license)

Please describe proprietary interest:

☐ C. I have significant financial interest vested in this proposal (remuneration that exceeds P250,000.00 each year or equity interest in the form of stock, stock options or other ownership interests).

Please describe financial interest:

A.3 Checklist A - Human Participants

 De La Salle University	Research Ethics Review Committee Research Ethics Office, 3F Henry Sy Sr. Hall De La Salle University Manila 2401 Taft Avenue, Manila 1004, Philippines REO@dlsu.edu.ph (632) 524-4611 loc. 513	SOP No.: 2
		Form No.: 2.03
		Version No.: 1
		Effectivity Date: July 2016

DE LA SALLE UNIVERSITY

Checklist A

Research Ethics Checklist for Investigations involving Human Participants

*This checklist must be completed **AFTER** the De La Salle University Code of Research Ethics and Guide to Responsible Conduct of Research has been read and **BEFORE** gathering data. The University Code of Research Ethics is available at http://www.dlsu.edu.ph/offices/urco/forms/URCO-Code-of-Research-Ethics_August2011.pdf*

NOTE: *This checklist is completed after the research proponent fills out the General Checklist Form.*

Only answer this Checklist if you answered YES on question 1 of the General Checklist.

Researcher Details	
Students	BERNANDINO , Adrian Rafael DC. BORROMEIO , Anton Miguel G. GALAN , Russell Emmanuel G. GANAYO , Reinman Geoffrey A.
Thesis Adviser	DEL GALLEG0 , Neil
Department/College	Department of Software Technology / College of Computer Studies
Proposed Title of the Research	A New Rendering and Plume Simulation Framework for Visualizing Taal Volcano Eruption Plumes
Term(s) and academic year in which research project is to be undertaken	Term 3 of AY 25-26 to Term 3 of AY 26-27

Provide a brief description of the data collection procedure to be undertaken in the research:

Data collection consists of a two-pronged approach: a technical hardware benchmarking phase and a human usability testing phase. First, the research team will execute automated, software-driven rendering loops comparing the legacy OpenGL engine against experimental low-level explicit graphics API implementations (e.g., Vulkan or DirectX 12) under identical hardware environments to capture frames per second (FPS), frame render times, and CPU/GPU utilization logs. Second, targeted human usability tests will be conducted following a tiered recruitment strategy: PHIVOLCS personnel will be prioritized as the primary participant group given their domain expertise in volcanology; should this pool be

 De La Salle University	Research Ethics Review Committee Research Ethics Office, 3F Henry Sy Sr. Hall De La Salle University Manila 2401 Taft Avenue, Manila 1004, Philippines REO@dlsu.edu.ph (632) 524-4611 loc. 513	SOP No.: 2
		Form No.: 2.03
		Version No.: 1
		Effectivity Date: July 2016

insufficient, recruitment will expand, in order, to individuals who have witnessed a Taal Volcano eruption firsthand, then to volcano/geohazard simulation enthusiasts, and lastly to computer science students with a background in rendering or volcano-related coursework. Participants will interact with the multi-API 3D plume simulator configurations. After completing standard interaction tasks, participants will complete an anonymous, self-administered questionnaire using the System Usability Scale (SUS) to evaluate application responsiveness, stability, and visual parity.

The following should be attached to the checklist:

- A copy of the informed consent form to be used in the study.
- A copy of the instrument/tool that will be administered to the participants.
- If applicable, a copy of the letter seeking permission to collect data from participants who are under the supervision of an agency, institution, department, or office.
- If applicable, a copy of the parental consent form for participants below 18 years old.
- If applicable, a copy of the letter seeking permission to collect data from participants who are under the supervision of an agency, institution, department, or office

The following items refer to important ethical considerations in the conduct of research with human participants. Provide a check for the appropriate answer to each question.

Source of data

Please check all that apply:

✓	1. New data will be collected from human participants If you checked this item, how will the new data be gathered? Please check all that apply. After answering this question, please proceed to page 3
	☐ Experimental Procedures/Intervention/ Treatments
	☐ Focus Group
	☐ Personal Interviews
	☒ Self-administered Questionnaire
	☐ Researcher-administered Questionnaire
	☐ Internet survey
	☐ Observation
	☐ Telephone survey
	☐ Others, please specify:

 De La Salle University	Research Ethics Review Committee Research Ethics Office, 3F Henry Sy Sr. Hall De La Salle University Manila 2401 Taft Avenue, Manila 1004, Philippines REO@dlsu.edu.ph (632) 524-4611 loc. 513	SOP No.: 2
		Form No.: 2.03
		Version No.: 1
		Effectivity Date: July 2016

	2. Pre-existing data from human participants, i.e., from a dataset you checked this item, please proceed to page 7
--	---

options are checked (both new data and pre-existing data), **answer all of the questions in this document.**

Only answer if new data will be collected (item 1 above)	
Sampling Details	
Number of Participants/Subjects	10 to 20 voluntary participants, recruited according to a tiered priority order: PHIVOLCS personnel as the primary target group, followed by individuals who witnessed a Taal Volcano eruption firsthand, then volcano/geohazard simulation enthusiasts, and lastly computer science students with a background in rendering or volcano-related coursework, should earlier priority groups be insufficient to reach the target sample size.
Location where the participants will be recruited/ where subjects will be obtained?	For PHIVOLCS personnel, recruitment will be initiated via email correspondence with the agency; if permission is granted, in-person coordination will also be done. For the remaining participant categories (Taal Volcano eruption witnesses, volcano/geohazard simulation enthusiasts, and computer science students or developers with a relevant background in 3D graphics or software development), recruitment will be conducted through a general online survey containing screening questions to confirm which category, if any, each respondent qualifies under.
How long will the data collection take place?	2 to 4 weeks during the system's evaluation timeframe.
Who will perform the data collection?	The primary student proponents (Bernandino, Borromeo, Galan, and Ganayo) under the adviser's supervision.
Location(s) where data collection will take place	DLSU GameLab, Gokongwei Building, Room 411 (G411).

 De La Salle University	Research Ethics Review Committee Research Ethics Office, 3F Henry Sy Sr. Hall De La Salle University Manila 2401 Taft Avenue, Manila 1004, Philippines REO@dlsu.edu.ph (632) 524-4611 loc. 513	SOP No.: 2
		Form No.: 2.03
		Version No.: 1
		Effectivity Date: July 2016

What procedures will be employed to ensure voluntary consent from participants?	Prior to the evaluation session, participants will receive a printed or digital Informed Consent Form (ICF) explicitly stating that participation is entirely voluntary, they may exit the software at any stage, and can choose to withdraw from the evaluation without academic or personal penalty.
Data Retention	
How long will data with participant identifiers be kept after the publication of the first paper from the project?	0 days (No personal identifiers will be collected from the beginning; all records remain fully anonymous).
How long will anonymized data be kept after the publication of the first paper from the project?	4 to 5 years following final publication or degree conferral to allow for reproducibility audits, securely saved in password-protected cloud environments before permanent deletion.
Procedure for Informed Consent	
How will informed consent be recorded? (check all that applies) Reminder: please attach informed consent that will be used in the study	☒ Written Consent ☐ Audio-recorded Consent ☒ Online/Email recorded Consent ☐ Others, please specify:

If you will not obtain a recorded informed consent, answer the questions that follow:

Why does the waiver of informed consent not pose a threat to the welfare and rights of the participants?

Why is recording an informed consent not practical for the proposed study?

	Yes	No	Not Applicable
1. Will the research involve students who will be receiving course credits for their participation?		✓	

 De La Salle University	Research Ethics Review Committee Research Ethics Office, 3F Henry Sy Sr. Hall De La Salle University Manila 2401 Taft Avenue, Manila 1004, Philippines REO@dlsu.edu.ph (632) 524-4611 loc. 513	SOP No.: 2
		Form No.: 2.03
		Version No.: 1
		Effectivity Date: July 2016

If YES, please attach a copy of the consent form and a summary of the debriefing process that will help participants understand how their participation in the research has provided a relevant learning experience to the crediting course.			
2. Does the study involve participants below 18 years old or those who are unable to give their informed consent? If YES, please attach a copy of the parental consent form.		✓	
3. Is there a possibility that the research can induce physical and/or psychological harm to the participants? Will they experience pain or some discomfort as a result from their participation in the research? If YES, please attach an acceptable argument that outlines the benefits of doing the research and how they outweigh the cost of harming the participants.		✓	
4. Will the participants be deliberately falsely informed or made unaware that they are being observed? Will they be misled in a way that they will possibly object to or show unease when told of the real purpose of the study? YES, please attach an acceptable argument that outlines the benefits of doing the research and how they outweigh the cost of harming the participants.		✓	
5. Will the research involve the discussion of, or questions on, sensitive topics (e.g. sexual activity, substance abuse, or mental health)? If YES, please make sure that the informed consent form explicitly states that sensitive questions will be posed and that you will safeguard the anonymity of the participants and ensure confidentiality. Please attach a copy of your informed consent form and your instrument.		✓	

 De La Salle University	Research Ethics Review Committee Research Ethics Office, 3F Henry Sy Sr. Hall De La Salle University Manila 2401 Taft Avenue, Manila 1004, Philippines REO@dlsu.edu.ph (632) 524-4611 loc. 513	SOP No.: 2
		Form No.: 2.03
		Version No.: 1
		Effectivity Date: July 2016

	Yes	No	Not Applicable
6. Will the research involve the administration of drugs, or other substances to the participants? YES, please attach an acceptable argument that outlines the benefits of doing the research and how they outweigh the cost of harming the participants. Please also attach a description of the procedure that will ensure that the participants will be brought back to their physical and psychological states prior to their participation in the research.		✓	
7. Will biological samples (e.g. blood, saliva, urine) be obtained from the participants? If YES, will this involve invasive procedures? Please attach a description of these procedures.		✓	
8. Will genetic materials be obtained from the biological samples? If YES, please attach a description of the procedures that will ensure confidentiality. Please attach the informed consent form.		✓	
9. Will financial inducements (other than reasonable expenses, like transportation or meal allowances) be offered to the participants for their participation in their research? If YES, the researcher(s) should be mindful of how the inducements can influence the participants' responses or behaviors during the research. Indicate the financial inducements offered to the participants: 		✓	
10. Is there a possibility for groups or communities to be harmed by the dissemination of the research findings?		✓	

 De La Salle University	Research Ethics Review Committee Research Ethics Office, 3F Henry Sy Sr. Hall De La Salle University Manila 2401 Taft Avenue, Manila 1004, Philippines REO@dlsu.edu.ph (632) 524-4611 loc. 513	SOP No.: 2
		Form No.: 2.03
		Version No.: 1
		Effectivity Date: July 2016

If YES, please attach a description of procedures to ensure the anonymity and confidentiality of the research findings.			
11. Will the results of this study have a commercial value? If yes, do you intend to apply for a patent for the output of this research? Please check: _____ Yes _____ No		✓	

FOR PROPONENTS WHO WILL GATHER NEW DATA ONLY, PLEASE STOP ANSWERING.

Use of Pre-existing Data collected from Human Participants		
Indicate the dataset from which the data for the study will be sourced		
Is the data publicly available, i.e., the access to which does not necessitate an approval process?		Yes Please indicate where the dataset is available:
		No Please indicate/attach the approval authority for access:
Was the original dataset originally collected for the present study's purpose?		Yes Please attach the Consent Form used in the original study.
		No Please attach the Information Collection Statement (i.e., the statement given to informants providing them with the rationale for the collection of specific information).
Does the original data set contain sensitive data, that is information that an individual would not likely want to be disclosed publicly, e.g., data on sexual activities, substance use?		Yes Please describe the type of sensitive data to be used in the present research:
		No
		No

 De La Salle University	Research Ethics Review Committee Research Ethics Office, 3F Henry Sy Sr. Hall De La Salle University Manila 2401 Taft Avenue, Manila 1004, Philippines REO@dlsu.edu.ph (632) 524-4611 loc. 513	SOP No.: 2
		Form No.: 2.03
		Version No.: 1
		Effectivity Date: July 2016

Does the original dataset have personal identifiers?		<i>(This means that neither the researcher nor the participant provided any personal identifiers)</i>
		Yes, specifically: _____ Direct (<i>i.e., the participant provided personal details like name and address</i>) _____ Indirect (<i>i.e., the participant was given a respondent code to make the participant identifiable</i>)
Will new data be collected and analyzed along with data from the existing dataset?		Yes Please answer questions on page 3-5.
		No

A.4 Informed Consent Form

A NEW RENDERING AND PLUME SIMULATION FRAMEWORK FOR VISUALIZING TAAI VOLCANO ERUPTION PLUMES

Informed Consent Form

This consent form is written for computer science students and developers with a relevant background in 3D graphics or software development.

Principal Investigators: Adrian Rafael Bernandino <adrian_rafael_bernandino@dlsu.edu.ph>, Anton Miguel Borromeo <anton_borromeo@dlsu.edu.ph>, Russell Emmanuel Galan <russell_galan@dlsu.edu.ph>, Reinman Geoffrey Ganayo <reinman_ganayo@dlsu.edu.ph>

Thesis Adviser: Neil Patrick Del Gallego, Ph.D. <neil.delgallego@dlsu.edu.ph>

Department/Institution: Department of Software Technology, College of Computer Studies, De La Salle University

Name of Research Project: A New Rendering and Plume Simulation Framework for Visualizing Taal Volcano Eruption Plumes

You may talk to anyone you feel comfortable talking with about the research, and you can take time to reflect on whether you want to participate. If you do not understand any of the words or concepts, we will take the time to explain them as we go, and you can ask questions at any time.

PURPOSE OF THE STUDY

You are being invited to take part in a research study. Before you decide to participate in this study, it is important that you understand why the research is being done and what your participation will involve. Please read the following information carefully, and feel free to ask the researcher if anything is unclear or if you need more information.

The purpose of the study is to evaluate a modernized 3D graphics rendering framework for an interactive Taal Volcano plume simulation (AnitoPlume V2). The study aims to assess whether the modernized system maintains an interactive and usable experience while evaluating rendering performance, graphics resource utilization, scalability, and visual fidelity. You are being invited to participate because you have a relevant background in computer science, software development, or 3D graphics, which is necessary to navigate and evaluate the simulator's technical features effectively.

STUDY PROCEDURES

If you choose to participate, you will be asked to complete the following procedures in chronological order:

- **Screening:** You will complete a short pre-screening form confirming that you meet the

inclusion criteria, such as being a collegiate CS student, graphics programmer, or software developer with prior exposure to real-time graphics or game development.

- **Orientation:** You will be given a brief, neutral walkthrough of the interactive Taal Volcano Plume simulator's controls and features.
- **Usability Task:** You will use the software to recreate specific Taal Volcano eruption plumes following a set of predefined scenarios.
- **Post-Session Questionnaire:** You will complete an anonymous, self-administered questionnaire using the System Usability Scale (SUS) to evaluate application responsiveness, stability, and visual parity.
- **Feedback:** You will participate in a closing interview to gather your suggestions for improving the interface and overall performance.

DURATION

The usability testing phase for the entire study will take place over 2 to 4 weeks. Your individual testing session will be kept brief to prevent eye strain and visual fatigue.

VOLUNTARY PARTICIPATION

Your participation in this study is voluntary. It is up to you whether or not you decide to participate. If you decide to participate, you will be asked to sign this consent form. After you sign this consent form, you are still free to withdraw at any time and without giving a reason. Withdrawing from this study will not affect the relationship you have, if any, with the researcher. You may exit the software at any stage and can choose to withdraw from the evaluation without academic or personal penalty. If you withdraw from the study before data collection is completed, your data will be destroyed.

RISKS

There are minimal risks associated with this study.

- **Physical:** You may experience minor eye strain or visual fatigue from interacting with the 3D graphics simulator.
- **Technical:** You may experience frustration if you encounter bugs or usability issues with the experimental prototype.

To mitigate these risks, researchers will provide immediate technical support for major difficulties. You are encouraged to request breaks at any time if you feel tired or frustrated, and the testing sessions will be kept brief.

You may decline to answer any or all questions and you may withdraw your participation at any time if you choose.

BENEFITS

There will be no direct financial benefit to your participation in the study. However, we hope that the information from this study will help establish a scalable rendering foundation that could potentially serve as one of the first accessible 3D visualization tools for volcanic hazard communication in the Philippines.

CONFIDENTIALITY

Your responses in this research will be anonymous. Every effort will be made by the researcher to preserve your confidentiality, including the following:

- All performance logs and usability questionnaire responses will be fully anonymized using alphanumeric identifiers.
- No personally identifiable information (PII) such as names, student numbers, or email addresses will be collected or tied to your responses.
- Survey sheets and digital spreadsheets will be secured in a password-protected cloud storage repository accessible only to the primary student researchers.
- Anonymized data will be kept for 4 to 5 years following final publication or degree conferral to allow for reproducibility audits before permanent deletion.
- Research results for publication will not include any personal identifiers of the research participant, unless explicit consent was given.
- Research participant will be notified and a separate consent will be asked should any collected data be used beyond the scope of this consent.

CONTACT INFORMATION

If you have any questions at any time about this study, or if you experience any non-normative sensations because of participation, or to exercise your data privacy rights, you may contact the researcher whose contact information is on the first page. If you have any questions regarding your rights as a research participant, or if problems arise which you do not feel you can discuss with the Principal Investigator, please contact the Chair of the DLSU Research Ethics Review Committee at chairrerc@dlsu.edu.ph (632) 524-4611 local 513.

CONSENT

This section is mandatory I have read the provided information, or it has been read to me. I have had the opportunity to ask questions about it and any questions I have asked have been answered to my satisfaction. I understand that I will be given a copy of this form, and the researcher will keep another copy on file. I consent voluntarily to be a participant in this study.

Print Name of Participant _____

Signature of Participant _____

Date _____

Day/month/year

This section is mandatory

Print Name of Impartial Witness _____

Signature of Impartial Witness _____

Date _____

Day/month/year

This section is mandatory

Print Name of Researcher _____

Signature of Researcher _____

Date _____

Day/month/year

If the participant is illiterate¹

I have witnessed the accurate reading of the consent form to the potential participant, and the individual has had the opportunity to ask questions. I confirm that the individual has given consent freely.

Print name of witness _____

Signature of witness _____

Date _____

Day/month/year

¹ A literate witness must sign (if possible, this person should be selected by the participant and should have no connection to the research team).

References

- Amador Herrera, J. A., Klein, J., Liu, D., Pałubicki, W., Pirk, S., & Michels, D. L. (2024). Cyclogenesis: Simulating hurricanes and tornadoes. *ACM Transactions on Graphics (TOG)*, 43(4), 1–16.
- Asgary, A., Bonadonna, C., & Frischknecht, C. (2019). Simulation and visualization of volcanic phenomena using microsoft hololens: case of vulcano island (italy). *IEEE transactions on engineering management*, 67(3), 545–553.
- Asplund, J. (2020). *Directx 12: Performance comparison between single- and multithreaded rendering when culling multiple lights* (Unpublished master’s thesis). Jönköping University, School of Engineering.
- Balangué-Tarriela, M., Lagmay, A., Sarmiento, D., Vasquez, J., Baldago, M., Ybañez, R., ... others (2022). Analysis of the 2020 taal volcano tephra fall deposits from crowdsourced information and field data. *Bulletin of volcanology*, 84(3), 35.
- Baldur Karlsson. (2023). *Renderdoc: Graphics debugging tool*. Retrieved from <https://renderdoc.org/> (Version 1.x; open-source frame capture and analysis tool for Vulkan, DirectX, OpenGL, and other graphics APIs)
- Bartosz Taudul. (n.d.). *Tracy: A real time, nanosecond resolution, remote telemetry frame profiler*. Retrieved from <https://github.com/wolfpld/tracy> (Open-source, cross-platform profiler supporting GPU zones, memory tracking, and live profilers)
- Biass, S., Bonadonna, C., Connor, L., & Connor, C. (2016). Tephraprob: a matlab package for probabilistic hazard assessments of tephra fallout. *Journal of Applied Volcanology*, 5(1), 10.
- Blender Foundation. (2025). *Blender - the free and open source 3D creation software*. <https://www.blender.org/>. (Accessed: 2026-07-10)
- Bonadonna, C. (2006). Probabilistic modelling of tephra dispersion.
- Bossard, A. (2019). High-performance graphics with directx in racket. *Information Engineering Express*, 5(2), 83–98. doi: 10.52731/iee.v5.i2.368
- Brackbill, J. U., & Ruppel, H. M. (1986). FLIP: A method for adaptively zoned, particle-in-cell calculations of fluid flows in two dimensions. *Journal of Computational Physics*, 65(2), 314–343. Retrieved from <https://doi.org/>

- 10.1016/0021-9991(86)90211-1 doi: 10.1016/0021-9991(86)90211-1
- Braun, B., Bender, J., & Thuerey, N. (2025). Adaptive phase-field-flip for very large scale two-phase fluid simulation. *ACM Transactions on Graphics (TOG)*, 44(4), 1–23.
- Capodiecici, N., Cavicchioli, R., & Marongiu, A. (2022). A taxonomy of modern gpgpu programming methods: On the benefits of a unified specification. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, 41, 1649–1662. doi: 10.1109/tcad.2021.3082863
- Cheung, D. (2024). *Applications and limitations of vulkan-layer-based instrumentation* (Unpublished master’s thesis). University of Alberta (Department of Computing Science).
- Connor, L. J., & Connor, C. B. (2006). Inversion is the key to dispersion: understanding eruption dynamics by inverting tephra fallout.
- Daguno-Bersamina, K. (2023). *Volcanic smog detected in Taal, health advisory issued*. Retrieved from <https://www.philstar.com/headlines/2023/09/22/2298189/volcanic-smog-detected-taal-health-advisory-issued> (Online News Article. Published September 22, 2023)
- Delage, E., van Wyk de Vries, B., Philippe, M., Conway, S., Morino, C., Manrique Llerena, N., ... Kristinn Helgason, J. (2021). Visualising and experiencing geological flows in virtual reality. In *Egu general assembly conference abstracts* (pp. EGU21–8801).
- Del Castillo, M., Paraiso, P., Vicente, M., Jamero, M., & Narisma, G. (2020). Impacts of taal volcano phreatic eruption (12 january 2020) on the environment and population: satellite-based observations compared with historical records. *Manila Observatory*.
- Del Gallego, N. P., Torre, O. D., Arquillo, C. M., Cordero, V. V., Sales, J. L., & Yarte, S. L. (2026). Interactive 3d simulation of taal volcano eruption plumes. *Applied Computing and Geosciences*, 100331.
- Driedger, C. L., Ramsey, D. W., Scott, W. E., Faust, L. M., Bard, J. A., & Wold, P. (2024). Following the tug of the audience from complex to simplified hazards maps at cascade range volcanoes. *Journal of Applied Volcanology*, 13(1), 4.
- Esposti Ongaro, T., Cerminara, M., Charbonnier, S. J., Lube, G., & Valentine, G. A. (2020). A framework for validation and benchmarking of pyroclastic current models. *Bulletin of Volcanology*, 82, 59. doi: 10.1007/s00445-020-01386-4
- Fearnley, C., & Beaven, S. (2018). Volcano alert level systems: managing the challenges of effective volcanic crisis communication. *Bulletin of Volcanology*, 80(5), 46.
- Fedkiw, R. (2003). Simulating natural phenomena. In *Geometric level set methods in imaging, vision, and graphics* (pp. 461–479). Springer.
- Flores, H. (2019). *Taal volcano showing signs of activity?* Re-

- trieved from <https://www.philstar.com/nation/2019/12/02/1973465/taal-volcano-showing-signs-activity> (Online News Article.)
- Folch, A., Costa, A., & Macedonio, G. (2009). Fall3d: A computational model for transport and deposition of volcanic ash. *Computers & Geosciences*, 35(6), 1334–1342.
- Folch, A., Costa, A., & Macedonio, G. (2016). Fplume-1.0: An integral volcanic plume model accounting for ash aggregation. *Geoscientific Model Development*, 9(1), 431–450. Retrieved from <https://gmd.copernicus.org/articles/9/431/2016/> doi: 10.5194/gmd-9-431-2016
- Folch, A., Mingari, L., Gutierrez, N., Hanzich, M., Macedonio, G., & Costa, A. (2020). Fall3d-8.0: a computational model for atmospheric transport and deposition of particles, aerosols and radionuclides—part 1: Model physics and numerics. *Geoscientific Model Development*, 13, 1431–1458. doi: 10.5194/gmd-13-1431-2020
- George, D. L., & Iverson, R. M. (2014, 10). A depth-averaged debris-flow model that includes the effects of evolving dilatancy. II. numerical predictions and experimental tests. *Proceedings of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 470(2170), 20130820. Retrieved from <https://doi.org/10.1098/rspa.2013.0820> doi: 10.1098/rspa.2013.0820
- Guerrero Tello, J. F., Coltelli, M., Marsella, M., Celauro, A., & Palenzuela Baena, J. A. (2022). Convolutional neural network algorithms for semantic segmentation of volcanic ash plumes using visible camera imagery. *Remote Sensing*, 14(18), 4477. doi: 10.3390/rs14184477
- Hansen, K., Barker, S., Illsley-Kemp, F., Anslow, C., Mueller, C., & Leonard, G. (2025). Volcanovr: A virtual reality environment for volcanic data visualisation and communication. *Volcanica*, 8(1), 175–187.
- Harris, M. (2004). Mapping computational concepts to gpus. *Electronic Engineering Times*(1351), 52-.
- Ioannidis, C., & Boutsis, A.-M. (2020). Multithreaded rendering for cross-platform 3D visualization based on Vulkan API. In *International archives of the photogrammetry, remote sensing and spatial information sciences* (Vol. XLIV-4/W1-2020, pp. 57–62). Göttingen: Copernicus GmbH.
- Iverson, R. M., & George, D. L. (2014). A depth-averaged debris-flow model that includes the effects of evolving dilatancy. i. physical basis. *Proceedings of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 470(2170), 20130819. doi: 10.1098/rspa.2013.0819
- Jänicke, H., & Scheuermann, G. (2010). Measuring complexity in lagrangian and eulerian flow descriptions. In *Computer graphics forum* (Vol. 29, pp. 1783–1794).
- Kavanagh, J. L., Engwell, S. L., & Martin, S. A. (2018). A review of laboratory and numerical modelling in volcanology. *Solid Earth*, 9(2), 531–571.

- Kerlow, I., Pedreros, G., & Albert, H. (2020). Earth girl volcano: characterizing and conveying volcanic hazard complexity in an interactive casual game of disaster preparedness and response. *Geoscience Communication*, 3(2), 343–364.
- Khronos Group. (2024). Vulkan specification (Version 1.4 ed.) [Computer software manual]. Retrieved from <https://registry.khronos.org/vulkan/specs/latest/html/vkspec.html>
- Kinkor, J. (2022). Easy vulkan. In *Proceedings of the 27th student eeict 2022*. Brno University of Technology, Faculty of Information Technology.
- Lastic, M., Rohmer, D., Cordonnier, G., Jaupart, C., Neyret, F., & Cani, M.-P. (2022). Interactive simulation of plume and pyroclastic volcanic ejections. *Proceedings of the ACM on Computer Graphics and Interactive Techniques*, 5(1), 1–15.
- Li, W. (2023). Synthesizing procedural landscape for volcanic eruption simulation in virtual reality. In *2023 7th international symposium on multidisciplinary studies and innovative technologies (ismsit)* (pp. 1–5).
- Luebke, D. (2008). Cuda: Scalable parallel programming for high-performance scientific computing. In *2008 5th ieee international symposium on biomedical imaging: From nano to macro* (p. 836–838). IEEE.
- Madera-Ramirez, F. A., Lopez-Martinez, J. L., Moo-Mena, F., & Gomez-Montalvo, J. (2022). Cpu-gpu buffer communication using compute shader to fill volumes with spheres. *The Journal of Supercomputing*, 78(5), 6448–6460.
- Mallari Jr., D. (2026). *Taal volcano activity increases with 120 quakes, 113 tremors*. Retrieved from <https://newsinfo.inquirer.net/2258129/taal-volcano-activity-increases-with-120-quakes-113-tremors> (Online News Article. Published July 6, 2026)
- Moyroud, N., & Portet, F. (2018). Introduction to qgis. *QGIS and generic tools*, 1, 1–17.
- Nabizadeh, M. S., Roy-Chowdhury, R., Yin, H., Ramamoorthi, R., & Chern, A. (2024). Fluid implicit particles on coadjoint orbits. *arXiv preprint arXiv:2406.01936*.
- Nave, R., Isaia, R., Vilardo, G., & Barclay, J. (2010). Re-assessing volcanic hazard maps for improving volcanic risk communication: application to stromboli island, italy. *Journal of Maps*, 6(1), 260–269.
- Nishita, T., & Dobashi, Y. (2001). Modeling and rendering of various natural phenomena consisting of particles. In *Proceedings. computer graphics international 2001* (pp. 149–156).
- Nowrouzezahrai, D. (2006). *Vortex based smoke simulation and control*.
- NVIDIA Corporation. (2023). *NVIDIA Nsight Graphics*. Retrieved from <https://developer.nvidia.com/nsight-graphics> (Need drone-isolated profiler for NVIDIA GPUs providing hardware relume metrics, frame de-

- bugging, and show creation)
- The OpenGL graphics system: A specification (Version 4.6 (Core Profile) ed.) [Computer software manual]. (2022, May). Retrieved from <https://registry.khronos.org/OpenGL/specs/gl/glspec46.core.pdf>
- openTopography. (2025). *OpenTopography*. <https://opentopography.org/>. (Accessed: 2026-07-10)
- Owens, J. D., Houston, M., Luebke, D., Green, S., Stone, J. E., & Phillips, J. C. (2008). Gpu computing. *Proceedings of the IEEE*, 96(5), 879-899.
- Pardini, F., Barsotti, S., Bonadonna, C., de' Michieli Vitturi, M., Folch, A., Mastin, L., ... Prata, A. T. (2024). Dynamics, monitoring, and forecasting of tephra in the atmosphere. *Reviews of Geophysics*, 62(4), e2023RG000808. Retrieved from <https://doi.org/10.1029/2023RG000808> doi: 10.1029/2023RG000808
- Pfaff, T., Thuerey, N., & Gross, M. (2012). Lagrangian vortex sheets for animating fluids. *ACM Transactions on Graphics (TOG)*, 31(4). doi: 10.1145/2185520.2185608
- PHIVOLCS. (2024). *Vmepd: Eruption history*. Retrieved from <https://wovodat.phivolcs.dost.gov.ph/volcan/erupt-history?volcan=580&sdate=&edate=&btn-search=>
- Ping, L., Qi, S., Chen, W., Jie, G., Yanwen, G., & Wenzhe, S. (2026, March). Modern graphics APIs: Design principles, a use case, and new perspectives. *ZTE Communications*, 24(1), 97–106. Retrieved from https://www.zte.com.cn/content/dam/zte-site/res-www-zte-com-cn/mediare/magazine/publication/com_en/article/en202601/20260113.pdf
- Poland, M. P., & Anderson, K. R. (2020). Partly cloudy with a chance of lava flows: Forecasting volcanic eruptions in the twenty-first century. *Journal of Geophysical Research: Solid Earth*, 125(1), e2018JB016974.
- Preppernau, C. A., & Jenny, B. (2015). Three-dimensional versus conventional volcanic hazard maps. *Natural Hazards*, 78(2), 1329–1347.
- Pretorius, P. C., Gain, J., Lastic, M., Cordonnier, G., Chen, J., Rohmer, D., & Cani, M.-P. (2024). Volcanic skies: coupling explosive eruptions with atmospheric simulation to create consistent skiescapes. In *Computer graphics forum* (Vol. 43, p. e15034).
- Reeves, W. T. (1998). Particle systems—a technique for modeling a class of fuzzy objects. In *Seminal graphics: pioneering efforts that shaped the field* (pp. 203–220).
- Renzo, A. (2020). *List: Jan. 14 class suspensions due to taal eruption*. Online News Article. Retrieved from <https://newsinfo.inquirer.net/1212526/list-january-14-class-suspensions-due-to-taal-eruption>
- Rossant, C., & Rougier, N. P. (2021). High-performance interactive scientific visualization with datoviz via the vulkan low-level gpu api. *Authorea Preprints*.

doi: 10.22541/au.162129047.76547134/v1

- Roth, M., & Guritz, R. (1995). Visualization of volcanic ash clouds. *IEEE computer graphics and applications*, 15(4), 34–39.
- Saidi, M. S., Rismanian, M., Monjezi, M., Zendehbad, M., & Fatehiboroujeni, S. (2014). Comparison between lagrangian and eulerian approaches in predicting motion of micron-sized particles in laminar flows. *Atmospheric Environment*, 89, 199–206.
- Sanchez-Gonzalez, A., Godwin, J., Pfaff, T., Ying, R., Leskovec, J., & Battaglia, P. (2020). Learning to simulate complex physics with graph networks. In *International conference on machine learning* (pp. 8459–8468).
- Schütz, M., Kerbl, B., & Wimmer, M. (2021). Rendering point clouds with compute shaders and vertex order optimization. *Computer Graphics Forum*, 40(4), 115–126. (Eurographics Symposium on Rendering 2021) doi: 10.1111/cgf.14345
- Selle, A., Rasmussen, N., & Fedkiw, R. (2005). A vortex particle method for smoke, water and explosions. In *Acm siggraph 2005 papers* (pp. 910–914).
- Shiraeef, J. (2016). *An exploratory study of high performance graphics application programming interfaces* (Unpublished master’s thesis). University of Tennessee at Chattanooga. (UTC Digital Commons)
- Shiraeef, J. A. (2016). *An exploratory study of high performance graphics application programming interfaces* (Master’s thesis, University of Tennessee at Chattanooga). Retrieved from <https://scholar.utc.edu/theses/446/>
- Smithsonian Institution. (2026). *Taal*. Retrieved from <https://volcano.si.edu/volcano.cfm?vn=273070> (Online Website page)
- Stam, J. (2023). Stable fluids. In *Seminal graphics papers: Pushing the boundaries, volume 2* (pp. 779–786).
- Stein, A., Draxler, R., Rolph, G., Stunder, B., Cohen, M., Ngan, F., ... others (2015). *Noaa’s hysplit atmospheric transport and dispersion modeling system, b. am. meteorol. soc., 96, 2059–2077*.
- Stora, D., Agliati, P.-O., Cani, M.-P., Neyret, F., & Gascuel, J.-D. (1999). Animating lava flows. In *Graphics interface (gi’99) proceedings* (pp. 203–210).
- Swetha, G., Datla, R., & Vishnu, C. (2023, July). MS-VACSNet: A network for multi-scale volcanic ash cloud segmentation in remote sensing images. In *2023 18th international conference on machine vision and applications (mva)* (pp. 1–6). IEEE.
- TecAdmin Staff. (2025). *Cpu vs. gpu: Unpacking the core differences*. Retrieved from <https://tecadmin.net/cpu-vs-gpu-key-differences/> (Online Article Image.)
- TechArx Gaming Staff. (2016). *Understanding graphics apis*. Retrieved from <https://techarx.com/understanding-graphics-apis/> (Online Article Image.)
- Thompson, M. A., Lindsay, J. M., & Gaillard, J.-C. (2015). The influence of prob-

- abilistic volcanic hazard map properties on hazard communication. *Journal of Applied Volcanology*, 4(1), 6.
- Thompson, M. A., Lindsay, J. M., & Leonard, G. S. (2017). More than meets the eye: volcanic hazard map design and visual communication. In *Observing the volcano world: volcano crisis communication* (pp. 621–640). Springer.
- Tibaldi, A., Bonali, F. L., Vitello, F., Delage, E., Nomikou, P., Antoniou, V., ... Whitworth, M. (2020). Real world-based immersive virtual reality for research, teaching and communication in volcanology. *Bulletin of Volcanology*, 82(5), 38.
- Torrise, F., Corradino, C., Cariello, S., & Del Negro, C. (2024). Enhancing detection of volcanic ash clouds from space with convolutional neural networks. *Journal of Volcanology and Geothermal Research*, 448, 108046. doi: 10.1016/j.jvolgeores.2024.108046
- Trujillo-Vela, M. G., Ramos-Cañon, A. M., Escobar-Vargas, J. A., & Galindo-Torres, S. A. (2022). An overview of debris-flow mathematical modelling. *Earth-Science Reviews*, 232, 104135. Retrieved from <https://doi.org/10.1016/j.earscirev.2022.104135> doi: 10.1016/j.earscirev.2022.104135
- Unterguggenberger, J., Kerbl, B., & Wimmer, M. (2022). The road to vulkan: Teaching modern low-level apis in introductory graphics courses. In *Eurographics 2022 - education papers*. The Eurographics Association. doi: 10.2312/eged.20221043
- U.S. Geological Survey. (2018). *Standard forest metrics - USGS 3DEP LiDAR*. Retrieved from <https://eros.usgs.gov/doi-remote-sensing-activities/2018/usgs/standard-forest-metrics-usgs-3dep-lidar> (Online Article/Report.)
- Va, H., Choi, M.-H., & Hong, M. (2021). Real-time cloth simulation using compute shader in unity3d for ar/vr contents. *Applied Sciences*, 11(17), 8255. Retrieved from <https://doi.org/10.3390/app11178255> doi: 10.3390/app11178255
- Vimont, U., Gain, J., Lastic, M., Cordonnier, G., Abiodun, B., & Cani, M.-P. (2020). Interactive meso-scale simulation of skylines. In *Computer graphics forum* (Vol. 39, pp. 585–596).
- Wilkes, T. C., Perring, T. D., & McGonigle, A. J. (2022). Semantic segmentation of explosive volcanic plumes through deep learning. *Computers & Geosciences*, 168, 105216. doi: 10.1016/j.cageo.2022.105216
- Zhang, S., Kong, F., Li, C., Wang, C., & Qin, H. (2017). Hybrid modeling of multiphysical processes for particle-based volcano animation. *Computer Animation and Virtual Worlds*, 28(3-4), e1758.
- Zoraja, I., Bonkovic, M., Papic, V., & Sunderam, V. (2023). Vanityx: An agile 3d rendering platform supporting mixed reality. *Applied Sciences*, 13, 5468. Retrieved from <https://doi.org/10.3390/app13095468> doi: 10.3390/app13095468